

# Personalized cancer diagnosis

## 1. Business Problem

### 1.1. Description

Source: <https://www.kaggle.com/c/msk-redefining-cancer-treatment/>

Data: Memorial Sloan Kettering Cancer Center (MSKCC)

Download training\_variants.zip and training\_text.zip from Kaggle.

#### **Context:**

Source: <https://www.kaggle.com/c/msk-redefining-cancer-treatment/discussion/35336#198462>

#### **Problem statement :**

Classify the given genetic variations/mutations based on evidence from text-based clinical literature.

### 1.2. Source/Useful Links

Some articles and reference blogs about the problem statement

1. <https://www.forbes.com/sites/matthewherper/2017/06/03/a-new-cancer-drug-helped-almost-everyone-who-took-it-almost-heres-what-it-teaches-us/#2a44ee2f6b25>
2. <https://www.youtube.com/watch?v=UwbuW7oK8rk>
3. <https://www.youtube.com/watch?v=qxXRKVompl8>

### 1.3. Real-world/Business objectives and constraints.

- No low-latency requirement.
- Interpretability is important.
- Errors can be very costly.
- Probability of a data-point belonging to each class is needed.

## 2. Machine Learning Problem Formulation

### 2.1. Data

#### 2.1.1. Data Overview

- Source: <https://www.kaggle.com/c/msk-redefining-cancer-treatment/data>
- We have two data files: one contains the information about the genetic mutations and the other contains the clinical evidence (text) that human experts/pathologists use to classify the genetic mutations.
- Both these data files have a common column called ID
- Data file's information:
  - training\_variants (ID , Gene, Variations, Class)
  - training\_text (ID, Text)

#### 2.1.2. Example Data Point

## training\_variants

ID, Gene, Variation, Class  
0, FAM58A, Truncating Mutations, 1  
1, CBL, W802\*, 2  
2, CBL, Q249E, 2  
...

## training\_text

ID, Text  
0|Cyclin-dependent kinases (CDKs) regulate a variety of fundamental cellular processes. CDK10 stands out as one of the last orphan CDKs for which no activating cyclin has been identified and no kinase activity revealed. Previous work has shown that CDK10 silencing increases ETS2 (v-ets erythroblastosis virus E26 oncogene homolog 2)-driven activation of the MAPK pathway, which confers tamoxifen resistance to breast cancer cells. The precise mechanisms by which CDK10 modulates ETS2 activity, and more generally the functions of CDK10, remain elusive. Here we demonstrate that CDK10 is a cyclin-dependent kinase by identifying cyclin M as an activating cyclin. Cyclin M, an orphan cyclin, is the product of FAM58A, whose mutations cause STAR syndrome, a human developmental anomaly whose features include toe syndactyly, telecanthus, and anogenital and renal malformations. We show that STAR syndrome-associated cyclin M mutants are unable to interact with CDK10. Cyclin M silencing phenocopies CDK10 silencing in increasing c-Raf and in conferring tamoxifen resistance to breast cancer cells. CDK10/cyclin M phosphorylates ETS2 in vitro, and in cells it positively controls ETS2 degradation by the proteasome. ETS2 protein levels are increased in cells derived from a STAR patient, and this increase is attributable to decreased cyclin M levels. Altogether, our results reveal an additional regulatory mechanism for ETS2, which plays key roles in cancer and development. They also shed light on the molecular mechanisms underlying STAR syndrome. Cyclin-dependent kinases (CDKs) play a pivotal role in the control of a number of fundamental cellular processes (1). The human genome contains 21 genes encoding proteins that can be considered as members of the CDK family owing to their sequence similarity with bona fide CDKs, those known to be activated by cyclins (2). Although discovered almost 20 y ago (3, 4), CDK10 remains one of the two CDKs without an identified cyclin partner. This knowledge gap has largely impeded the exploration of its biological functions. CDK10 can act as a positive cell cycle regulator in some cells (5, 6) or as a tumor suppressor in others (7, 8). CDK10 interacts with the ETS2 (v-ets erythroblastosis virus E26 oncogene homolog 2) transcription factor and inhibits its transcriptional activity through an unknown mechanism (9). CDK10 knockdown derepresses ETS2, which increases the expression of the c-Raf protein kinase, activates the MAPK pathway, and induces resistance of MCF7 cells to tamoxifen (6). ...

## 2.2. Mapping the real-world problem to an ML problem

### 2.2.1. Type of Machine Learning Problem

There are nine different classes a genetic mutation can be classified into => Multi class classification problem

### 2.2.2. Performance Metric

Source: <https://www.kaggle.com/c/msk-redefining-cancer-treatment#evaluation>

Metric(s):

- Multi class log-loss
- Confusion matrix

### 2.2.3. Machine Learning Objectives and Constraints

Objective: Predict the probability of each data-point belonging to each of the nine classes.

Constraints:

- Interpretability
- Class probabilities are needed.
- Penalize the errors in class probabilities => Metric is Log-loss.
- No Latency constraints.

## 2.3. Train, CV and Test Datasets

Split the dataset randomly into three parts train, cross validation and test with 64%,16%, 20% of data respectively

## 3. Exploratory Data Analysis

In [1]:

```
import pandas as pd
import matplotlib.pyplot as plt
import re
import time
import warnings
import numpy as np
from nltk.corpus import stopwords
from sklearn.decomposition import TruncatedSVD
from sklearn.preprocessing import normalize
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.manifold import TSNE
import seaborn as sns
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import confusion_matrix
from sklearn.metrics.classification import accuracy_score, log_loss
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.linear_model import SGDClassifier
from imblearn.over_sampling import SMOTE
from collections import Counter
from scipy.sparse import hstack
from sklearn.multiclass import OneVsRestClassifier
#from sklearn import OneVsRestClassifier
from sklearn.svm import SVC
from sklearn.model_selection import StratifiedKFold
from collections import Counter, defaultdict
from sklearn.calibration import CalibratedClassifierCV
from sklearn.naive_bayes import MultinomialNB
from sklearn.naive_bayes import GaussianNB
from sklearn.model_selection import train_test_split
from sklearn.model_selection import GridSearchCV
import math
from sklearn.metrics import normalized_mutual_info_score
from sklearn.ensemble import RandomForestClassifier
warnings.filterwarnings("ignore")

from mlxtend.classifier import StackingClassifier

from sklearn import model_selection
from sklearn.linear_model import LogisticRegression
```

## 3.1. Reading Data

### 3.1.1. Reading Gene and Variation Data

In [2]:

```
data = pd.read_csv('training_variants')
print('Number of data points : ', data.shape[0])
print('Number of features : ', data.shape[1])
print('Features : ', data.columns.values)
data.head()
```

```
Number of data points : 3321
Number of features : 4
Features : ['ID' 'Gene' 'Variation' 'Class']
```

Out[2]:

	ID	Gene	Variation	Class

0	ID	Gene	Truncating Mutations	1	Class
1	1	CBL	W802*	2	
2	2	CBL	Q249E	2	
3	3	CBL	N454D	3	
4	4	CBL	L399V	4	

training/training\_variants is a comma separated file containing the description of the genetic mutations used for training.

Fields are

- **ID** : the id of the row used to link the mutation to the clinical evidence
- **Gene** : the gene where this genetic mutation is located
- **Variation** : the aminoacid change for this mutations
- **Class** : 1-9 the class this genetic mutation has been classified on

### 3.1.2. Reading Text Data

In [3]:

```
# note the separator in this file
data_text = pd.read_csv("training_text", sep="\\|\\|", engine="python", names=["ID", "TEXT"], skiprows=1)
print('Number of data points : ', data_text.shape[0])
print('Number of features : ', data_text.shape[1])
print('Features : ', data_text.columns.values)
data_text.head()
```

Number of data points : 3321

Number of features : 2

Features : ['ID' 'TEXT']

Out[3]:

	ID	TEXT
0	0	Cyclin-dependent kinases (CDKs) regulate a var...
1	1	Abstract Background Non-small cell lung canc...
2	2	Abstract Background Non-small cell lung canc...
3	3	Recent evidence has demonstrated that acquired...
4	4	Oncogenic mutations in the monomeric Casitas B...

### 3.1.3. Preprocessing of text

In [4]:

```
# loading stop words from nltk library
stop_words = set(stopwords.words(
    ('english')))

def nlp_preprocessing(total_text, index, column):
    if type(total_text) is not int:
        string = ""
        # replace every special char with space
        total_text = re.sub('[^a-zA-Z0-9\\n]', ' ', total_text)
        # replace multiple spaces with single space
        total_text = re.sub('\\s+', ' ', total_text)
        # converting all the chars into lower-case.
        total_text = total_text.lower()

        for word in total_text.split():
            # if the word is a not a stop word then retain that word from the data
            if not word in stop_words:
```

```

        string += word + " "

    data_text[column][index] = string

```

In [5]:

```

#text processing stage.
start_time = time.clock()
for index, row in data_text.iterrows():
    if type(row['TEXT']) is str:
        nlp_preprocessing(row['TEXT'], index, 'TEXT')
    else:
        print("there is no text description for id:",index)
print('Time took for preprocessing the text :',time.clock() - start_time, "seconds")

```

```

there is no text description for id: 1109
there is no text description for id: 1277
there is no text description for id: 1407
there is no text description for id: 1639
there is no text description for id: 2755
Time took for preprocessing the text : 379.25390385832304 seconds

```

In [6]:

```

#merging both gene_variations and text data based on ID
result = pd.merge(data, data_text,on='ID', how='left')
result.head()

```

Out[6]:

	ID	Gene	Variation	Class	TEXT
0	0	FAM58A	Truncating Mutations	1	cyclin dependent kinases cdks regulate variety...
1	1	CBL	W802*	2	abstract background non small cell lung cancer...
2	2	CBL	Q249E	2	abstract background non small cell lung cancer...
3	3	CBL	N454D	3	recent evidence demonstrated acquired uniparen...
4	4	CBL	L399V	4	oncogenic mutations monomeric casitas b lineag...

In [7]:

```

result[result.isnull().any(axis=1)]

```

Out[7]:

	ID	Gene	Variation	Class	TEXT
1109	1109	FANCA	S1088F	1	NaN
1277	1277	ARID5B	Truncating Mutations	1	NaN
1407	1407	FGFR3	K508M	6	NaN
1639	1639	FLT1	Amplification	6	NaN
2755	2755	BRAF	G596C	7	NaN

In [8]:

```

result.loc[result['TEXT'].isnull(),'TEXT'] = result['Gene'] + ' '+result['Variation']

```

In [9]:

```

result[result['ID']==1109]

```

Out[9]:

	ID	Gene	Variation	Class	TEXT
1109	1109	FANCA	S1088F	1	FANCA S1088F

### 3.1.4. Test, Train and Cross Validation Split

#### 3.1.4.1. Splitting data into train, test and cross validation (64:20:16)

In [10]:

```
y_true = result['Class'].values
result.Gene      = result.Gene.str.replace('\s+', '_')
result.Variation = result.Variation.str.replace('\s+', '_')

# split the data into test and train by maintaining same distribution of output variable 'y_true'
[stratify=y_true]
X_train, test_df, y_train, y_test = train_test_split(result, y_true, stratify=y_true, test_size=0.2)

# split the train data into train and cross validation by maintaining same distribution of output
variable 'y_train' [stratify=y_train]
train_df, cv_df, y_train, y_cv = train_test_split(X_train, y_train, stratify=y_train, test_size=0.2)
```

We split the data into train, test and cross validation data sets, preserving the ratio of class distribution in the original data set

In [11]:

```
print('Number of data points in train data:', train_df.shape[0])
print('Number of data points in test data:', test_df.shape[0])
print('Number of data points in cross validation data:', cv_df.shape[0])
```

```
Number of data points in train data: 2124
Number of data points in test data: 665
Number of data points in cross validation data: 532
```

#### 3.1.4.2. Distribution of y\_i's in Train, Test and Cross Validation datasets

In [12]:

```
# it returns a dict, keys as class labels and values as the number of data points in that class
train_class_distribution = train_df['Class'].value_counts().sortlevel()
test_class_distribution = test_df['Class'].value_counts().sortlevel()
cv_class_distribution = cv_df['Class'].value_counts().sortlevel()

my_colors = 'rgbkymc'
train_class_distribution.plot(kind='bar')
plt.xlabel('Class')
plt.ylabel('Data points per Class')
plt.title('Distribution of yi in train data')
plt.grid()
plt.show()


# ref: argsort https://docs.scipy.org/doc/numpy/reference/generated/numpy.argsort.html
# -(train_class_distribution.values): the minus sign will give us in decreasing order
sorted_yi = np.argsort(-train_class_distribution.values)
for i in sorted_yi:
    print('Number of data points in class', i+1, ':', train_class_distribution.values[i], '(', np.round(
        (train_class_distribution.values[i]/train_df.shape[0]*100), 3), '%)')

print('-'*80)
my_colors = 'rgbkymc'
test_class_distribution.plot(kind='bar')
plt.xlabel('Class')
plt.ylabel('Data points per Class')
plt.title('Distribution of yi in test data')
plt.grid()
plt.show()
```

```
# ref: argsort https://docs.scipy.org/doc/numpy/reference/generated/numpy.argsort.html
# -(train_class_distribution.values): the minus sign will give us in decreasing order
sorted_yi = np.argsort(-test_class_distribution.values)
for i in sorted_yi:
    print('Number of data points in class', i+1, ':', test_class_distribution.values[i], '(', np.round(
    ((test_class_distribution.values[i]/test_df.shape[0]*100), 3), '%)'))


print('-'*80)
my_colors = 'rgbkymc'
cv_class_distribution.plot(kind='bar')
plt.xlabel('Class')
plt.ylabel('Data points per Class')
plt.title('Distribution of yi in cross validation data')
plt.grid()
plt.show()

# ref: argsort https://docs.scipy.org/doc/numpy/reference/generated/numpy.argsort.html
# -(train_class_distribution.values): the minus sign will give us in decreasing order
sorted_yi = np.argsort(-train_class_distribution.values)
for i in sorted_yi:
    print('Number of data points in class', i+1, ':', cv_class_distribution.values[i], '(', np.round(
    ((cv_class_distribution.values[i]/cv_df.shape[0]*100), 3), '%)'))
```



Number of data points in class 7 : 609 ( 28.672 %)  
 Number of data points in class 4 : 439 ( 20.669 %)  
 Number of data points in class 1 : 363 ( 17.09 %)  
 Number of data points in class 2 : 289 ( 13.606 %)  
 Number of data points in class 6 : 176 ( 8.286 %)  
 Number of data points in class 5 : 155 ( 7.298 %)  
 Number of data points in class 3 : 57 ( 2.684 %)  
 Number of data points in class 9 : 24 ( 1.13 %)  
 Number of data points in class 8 : 12 ( 0.565 %)


---



Number of data points in class 7 : 191 ( 28.722 %)  
 Number of data points in class 4 : 137 ( 20.602 %)  
 Number of data points in class 1 : 114 ( 17.143 %)  
 Number of data points in class 2 : 91 ( 13.684 %)  
 Number of data points in class 6 : 55 ( 8.271 %)  
 Number of data points in class 5 : 48 ( 7.218 %)

Number of data points in class 3 : 18 ( 2.707 %)  
Number of data points in class 9 : 7 ( 1.053 %)  
Number of data points in class 8 : 4 ( 0.602 %)

---



Number of data points in class 7 : 153 ( 28.759 %)  
Number of data points in class 4 : 110 ( 20.677 %)  
Number of data points in class 1 : 91 ( 17.105 %)  
Number of data points in class 2 : 72 ( 13.534 %)  
Number of data points in class 6 : 44 ( 8.271 %)  
Number of data points in class 5 : 39 ( 7.331 %)  
Number of data points in class 3 : 14 ( 2.632 %)  
Number of data points in class 9 : 6 ( 1.128 %)  
Number of data points in class 8 : 3 ( 0.564 %)

## 3.2 Prediction using a 'Random' Model

In a 'Random' Model, we generate the NINE class probabilities randomly such that they sum to 1.

In [13]:

```
# This function plots the confusion matrices given  $y_i$ ,  $y_i$ -hat.
def plot_confusion_matrix(test_y, predict_y):
    C = confusion_matrix(test_y, predict_y)
    # C = 9,9 matrix, each cell (i,j) represents number of points of class i are predicted class j

    A = ((C.T) / (C.sum(axis=1))).T
    #divid each element of the confusion matrix with the sum of elements in that column

    # C = [[1, 2],
    #      [3, 4]]
    # C.T = [[1, 3],
    #        [2, 4]]
    # C.sum(axis = 1)  axis=0 corresponds to columns and axis=1 corresponds to rows in two
    dimensional array
    # C.sum(axis =1) = [[3, 7]]
    # ((C.T)/(C.sum(axis=1))) = [[1/3, 3/7]
    #                             [2/3, 4/7]]

    # ((C.T)/(C.sum(axis=1))).T = [[1/3, 2/3]
    #                               [3/7, 4/7]]
    # sum of row elements = 1

    B = (C/C.sum(axis=0))
    #divid each element of the confusion matrix with the sum of elements in that row
    # C = [[1, 2],
    #      [3, 4]]
    # C.sum(axis = 0)  axis=0 corresponds to columns and axis=1 corresponds to rows in two
    dimensional array
    # C.sum(axis =0) = [[4, 6]]
    # (C/C.sum(axis=0)) = [[1/4, 2/6],
    #                       [3/4, 4/6]]

    labels = [1,2,3,4,5,6,7,8,9]
    # representing A in heatmap format
    print("-"*20, "Confusion matrix", "-"*20)
```



```

plt.figure(figsize=(20,7))
sns.heatmap(C, annot=True, cmap="YlGnBu", fmt=".3f", xticklabels=labels, yticklabels=labels)
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.show()

print("-"*20, "Precision matrix (Column Sum=1)", "-"*20)
plt.figure(figsize=(20,7))
sns.heatmap(B, annot=True, cmap="YlGnBu", fmt=".3f", xticklabels=labels, yticklabels=labels)
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.show()

# representing B in heatmap format
print("-"*20, "Recall matrix (Row sum=1)", "-"*20)
plt.figure(figsize=(20,7))
sns.heatmap(A, annot=True, cmap="YlGnBu", fmt=".3f", xticklabels=labels, yticklabels=labels)
plt.xlabel('Predicted Class')
plt.ylabel('Original Class')
plt.show()

```

In [14]:

```

# we need to generate 9 numbers and the sum of numbers should be 1
# one solution is to generate 9 numbers and divide each of the numbers by their sum
# ref: https://stackoverflow.com/a/18662466/4084039
test_data_len = test_df.shape[0]
cv_data_len = cv_df.shape[0]

# we create a output array that has exactly same size as the CV data
cv_predicted_y = np.zeros((cv_data_len,9))
for i in range(cv_data_len):
    rand_probs = np.random.rand(1,9)
    cv_predicted_y[i] = ((rand_probs/sum(sum(rand_probs))))[0])
print("Log loss on Cross Validation Data using Random Model",log_loss(y_cv,cv_predicted_y, eps=1e-15))

# Test-Set error.
#we create a output array that has exactly same as the test data
test_predicted_y = np.zeros((test_data_len,9))
for i in range(test_data_len):
    rand_probs = np.random.rand(1,9)
    test_predicted_y[i] = ((rand_probs/sum(sum(rand_probs))))[0])
print("Log loss on Test Data using Random Model",log_loss(y_test,test_predicted_y, eps=1e-15))


predicted_y = np.argmax(test_predicted_y, axis=1)
plot_confusion_matrix(y_test, predicted_y+1)

```


Log loss on Cross Validation Data using Random Model 2.456814966189978

Log loss on Test Data using Random Model 2.479952691325209


----- Confusion matrix -----



Precision matrix (Column Sum=1)



Recall matrix (Row sum=1)



### 3.3 Univariate Analysis

In [15]:

```
# code for response coding with Laplace smoothing.
# alpha : used for laplace smoothing
# feature: ['gene', 'variation']
# df: ['train_df', 'test_df', 'cv_df']
# algorithm
# -----
# Consider all unique values and the number of occurrences of given feature in train data dataframe
# build a vector (1*9) , the first element = (number of times it occurred in class1 + 10*alpha / number of times it occurred in total data+90*alpha)
# gv_dict is like a look up table, for every gene it store a (1*9) representation of it
# for a value of feature in df:
# if it is in train data:
# we add the vector that was stored in 'gv_dict' look up table to 'gv_fea'
# if it is not there is train:
# we add [1/9, 1/9, 1/9, 1/9, 1/9, 1/9, 1/9, 1/9, 1/9] to 'gv_fea'
# return 'gv_fea'
# -----

# get_gv_fea_dict: Get Gene variation Feature Dict
def get_gv_fea_dict(alpha, feature, df):
    # value_count: it contains a dict like
    # print(train_df['Gene'].value_counts())
    # output:
```

```

# output.
# {BRCA1      174
#      TP53    106
#      EGFR    86
#      BRCA2   75
#      PTEN    69
#      KIT     61
#      BRAF    60
#      ERBB2   47
#      PDGFRA  46
#      ...}
# print(train_df['Variation'].value_counts())
# output:
# {
# Truncating_Mutations      63
# Deletion                  43
# Amplification              43
# Fusions                    22
# Overexpression             3
# E17K                       3
# Q61L                       3
# S222D                       2
# P130S                       2
# ...
# }
value_count = train_df[feature].value_counts()

# gv_dict : Gene Variation Dict, which contains the probability array for each gene/variation
gv_dict = dict()

# denominator will contain the number of time that particular feature occurred in whole data
for i, denominator in value_count.items():
    # vec will contain (p(yi==1/Gi) probability of gene/variation belongs to particular class
    # vec is 9 dimensional vector
    vec = []
    for k in range(1,10):
        # print(train_df.loc[(train_df['Class']==1) & (train_df['Gene']=='BRCA1')])
        #
        # ID      Gene      Variation      Class
        # 2470    2470    BRCA1      S1715C      1
        # 2486    2486    BRCA1      S1841R      1
        # 2614    2614    BRCA1      M1R        1
        # 2432    2432    BRCA1      L1657P      1
        # 2567    2567    BRCA1      T1685A      1
        # 2583    2583    BRCA1      E1660G      1
        # 2634    2634    BRCA1      W1718L      1
        # cls_cnt.shape[0] will return the number of rows

        cls_cnt = train_df.loc[(train_df['Class']==k) & (train_df[feature]==i)]

        # cls_cnt.shape[0] (numerator) will contain the number of time that particular feature occurred in whole data
        vec.append((cls_cnt.shape[0] + alpha*10)/ (denominator + 90*alpha))

    # we are adding the gene/variation to the dict as key and vec as value
    gv_dict[i]=vec
return gv_dict

# Get Gene variation feature
def get_gv_feature(alpha, feature, df):
    # print(gv_dict)
    # {'BRCA1': [0.20075757575757575, 0.03787878787878788, 0.06818181818181817,
    0.13636363636363635, 0.25, 0.19318181818181818, 0.03787878787878788, 0.03787878787878788,
    0.03787878787878788],
    #      'TP53': [0.32142857142857145, 0.061224489795918366, 0.061224489795918366,
    0.27040816326530615, 0.061224489795918366, 0.066326530612244902, 0.051020408163265307, 0.051020408
    163265307, 0.056122448979591837],
    #      'EGFR': [0.056818181818181816, 0.21590909090909091, 0.0625, 0.06818181818181817,
    0.06818181818181817, 0.0625, 0.34659090909090912, 0.0625, 0.056818181818181816],
    #      'BRCA2': [0.13333333333333333, 0.060606060606060608, 0.060606060606060608,
    0.07878787878787878, 0.1393939393939394, 0.34545454545454546, 0.060606060606060608,
    0.060606060606060608, 0.060606060606060608],
    #      'PTEN': [0.069182389937106917, 0.062893081761006289, 0.069182389937106917,
    0.46540880503144655, 0.075471698113207544, 0.062893081761006289, 0.069182389937106917, 0.062893081
    761006289, 0.062893081761006289],
    #      'KIT': [0.066225165562913912, 0.25165562913907286, 0.072847682119205295,
    0.072847682119205295, 0.066225165562913912, 0.066225165562913912, 0.27152317880794702,
    0.066225165562913912, 0.066225165562913912],
    #      'BRAF': [0.06666666666666666, 0.17888888888888888, 0.07222222222222222,

```

```

# BRAF : [0.0000000000000000, 0.1999999999999999, 0.0733333333333333,
0.0733333333333333, 0.0933333333333333, 0.0800000000000000, 0.2999999999999999,
0.0666666666666666, 0.0666666666666666],
# ...
# }
gv_dict = get_gv_fea_dict(alpha, feature, df)
# value_count is similar in get_gv_fea_dict
value_count = train_df[feature].value_counts()

# gv_fea: Gene_variation feature, it will contain the feature for each feature value in the data
gv_fea = []
# for every feature values in the given data frame we will check if it is there in the train
data then we will add the feature to gv_fea
# if not we will add [1/9,1/9,1/9,1/9,1/9,1/9,1/9,1/9,1/9,1/9] to gv_fea
for index, row in df.iterrows():
    if row[feature] in dict(value_count).keys():
        gv_fea.append(gv_dict[row[feature]])
    else:
        gv_fea.append([1/9,1/9,1/9,1/9,1/9,1/9,1/9,1/9,1/9,1/9])
# gv_fea.append([-1,-1,-1,-1,-1,-1,-1,-1,-1,-1])
return gv_fea

```

when we calculate the probability of a feature belongs to any particular class, we apply laplace smoothing

- $(\text{numerator} + 10 \cdot \alpha) / (\text{denominator} + 90 \cdot \alpha)$

### 3.2.1 Univariate Analysis on Gene Feature

**Q1.** Gene, What type of feature it is ?

**Ans.** Gene is a categorical variable

**Q2.** How many categories are there and How they are distributed?

In [16]:

```

unique_genes = train_df['Gene'].value_counts()
print('Number of Unique Genes :', unique_genes.shape[0])
# the top 10 genes that occurred most
print(unique_genes.head(10))

```

```

Number of Unique Genes : 226
BRCA1      176
TP53       112
EGFR        88
PTEN        86
BRCA2       77
BRAF        66
KIT         63
ALK         48
PIK3CA      42
ERBB2       40
Name: Gene, dtype: int64

```

In [17]:

```

print("Ans: There are", unique_genes.shape[0], "different categories of genes in the train data, and they are distributed as follows",)

```

Ans: There are 226 different categories of genes in the train data, and they are distributed as follows


In [18]:

```

s = sum(unique_genes.values);
h = unique_genes.values/s;
plt.plot(h, label="Histogram of Genes")
plt.xlabel('Index of a Gene')
plt.ylabel('Number of Occurances')
plt.legend()


```

```
plt.grid()
plt.show()
```



In [19]:

```
c = np.cumsum(h)
plt.plot(c, label='Cumulative distribution of Genes')
plt.grid()
plt.legend()
plt.show()
```



### Q3. How to featurize this Gene feature ?

**Ans.** there are two ways we can featurize this variable check out this video:

<https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/handling-categorical-and-numerical-features/>

1. One hot Encoding
2. Response coding

We will choose the appropriate featurization based on the ML model we use. For this problem of multi-class classification with categorical features, one-hot encoding is better for Logistic regression while response coding is better for Random Forests.

In [20]:

```
#response-coding of the Gene feature
# alpha is used for laplace smoothing
alpha = 1
# train gene feature
train_gene_feature_responseCoding = np.array(get_gv_feature(alpha, "Gene", train_df))
# test gene feature
test_gene_feature_responseCoding = np.array(get_gv_feature(alpha, "Gene", test_df))
# cross validation gene feature
cv_gene_feature_responseCoding = np.array(get_gv_feature(alpha, "Gene", cv_df))
```

In [21]:

```
print("train_gene_feature_responseCoding is converted feature using response coding method. The shape of gene feature:", train_gene_feature_responseCoding.shape)
```

train\_gene\_feature\_responseCoding is converted feature using response coding method. The shape of gene feature: (2124, 9)

In [22]:

```
# one-hot encoding of Gene feature.
gene_vectorizer = TfidfVectorizer(max_features=2000, ngram_range=(1,4))
train_gene_feature_onehotCoding = gene_vectorizer.fit_transform(train_df['Gene'])
test_gene_feature_onehotCoding = gene_vectorizer.transform(test_df['Gene'])
cv_gene_feature_onehotCoding = gene_vectorizer.transform(cv_df['Gene'])
```

In [23]:

```
train_df['Gene'].head()
```

Out[23]:

```
1402    FGFR3
1330    MLH1
761    ERBB2
389    TP53
2330    JAK2
Name: Gene, dtype: object
```

In [24]:

```
gene_vectorizer.get_feature_names()
```

Out[24]:

```
['abl1',
 'acvr1',
 'ago2',
 'akt1',
 'akt2',
 'akt3',
 'alk',
 'apc',
 'ar',
 'araf',
 'arid1a',
 'arid1b',
 'arid2',
 'arid5b',
 'asx1',
 'atm',
 'aurka',
 'aurkb',
 'axin1',
 'axl',
 'b2m',
 'bap1',
 'bcl10',
 'bcl2l1',
 'bcl2l11',
 'bcl2l111',
 'bcor',
 'braf',
 'brca1',
 'brca2',
 'brd4',
 'brip1',
 'btk',
 'card11',
 'carm1',
 'casp8',
 'cb1',
 'ccnd1',
 'ccnd2',
 'ccnd3',
 'ccne1',
 'cdh1',
 'cdk12',
 'cdk4',
```

'cdk6',  
'cdk8',  
'cdkn1a',  
'cdkn1b',  
'cdkn2a',  
'cdkn2b',  
'cdkn2c',  
'cebpa',  
'chek2',  
'cic',  
'crebbp',  
'ctcf',  
'ctnnb1',  
'ddr2',  
'dicer1',  
'dnmt3a',  
'dusp4',  
'egfr',  
'eiflax',  
'elf3',  
'ep300',  
'epas1',  
'epcam',  
'erbb2',  
'erbb3',  
'erbb4',  
'ercc2',  
'ercc3',  
'ercc4',  
'erg',  
'errfil',  
'esr1',  
'etv6',  
'ewsr1',  
'ezh2',  
'fam58a',  
'fanca',  
'fancc',  
'fat1',  
'fbxw7',  
'fgf19',  
'fgf3',  
'fgfr1',  
'fgfr2',  
'fgfr3',  
'fgfr4',  
'flt3',  
'foxa1',  
'foxl2',  
'foxp1',  
'gata3',  
'gli1',  
'gnas',  
'h3f3a',  
'hla',  
'hnf1a',  
'hras',  
'idh1',  
'idh2',  
'igflr',  
'ikbke',  
'jak1',  
'jak2',  
'kdm5a',  
'kdm5c',  
'kdm6a',  
'kdr',  
'keap1',  
'kit',  
'kmt2a',  
'kmt2c',  
'kmt2d',  
'knstrn',  
'kras',  
'lats1',  
'lats2',  
'map2k1',

'map2k2',  
'map2k4',  
'map3k1',  
'mapk1',  
'med12',  
'mef2b',  
'met',  
'mga',  
'mlh1',  
'mpl',  
'msh2',  
'msh6',  
'mtor',  
'myc',  
'myd88',  
'myod1',  
'ncor1',  
'nfl',  
'nf2',  
'nfe2l2',  
'nfkb1a',  
'nkx2',  
'notch1',  
'nras',  
'nsd1',  
'ntrk1',  
'ntrk2',  
'ntrk3',  
'nup93',  
'pax8',  
'pbrm1',  
'pdgfra',  
'pdgfrb',  
'pik3ca',  
'pik3cb',  
'pik3cd',  
'pik3r1',  
'pik3r2',  
'pim1',  
'pms2',  
'pole',  
'ppm1d',  
'ppp2r1a',  
'ppp6c',  
'prdm1',  
'ptch1',  
'pten',  
'ptpn11',  
'ptprd',  
'ptprt',  
'rac1',  
'rad21',  
'rad50',  
'rad51b',  
'rad51c',  
'rad54l',  
'raf1',  
'rasa1',  
'rb1',  
'rbm10',  
'ret',  
'rheb',  
'rhoa',  
'riCTOR',  
'rit1',  
'ros1',  
'rras2',  
'runx1',  
'rxra',  
'rybp',  
'sdhb',  
'setd2',  
'sf3b1',  
'shoc2',  
'shq1',  
'smad2',  
'smad3',



```
'smad4',
'smarca4',
'smo',
'sos1',
'sox9',
'spop',
'src',
'srsf2',
'stag2',
'stat3',
'stk11',
'tcf3',
'tcf7l2',
'tert',
'tet1',
'tet2',
'tgfbr1',
'tgfbr2',
'tmprss2',
'tp53',
'tp53bp1',
'tsc1',
'tsc2',
'u2af1',
'vegfa',
'vhl',
'whsc1',
'xpo1',
'xrcc2',
'yap1']
```

In [25]:

```
print("train_gene_feature_onehotCoding is converted feature using one-hot encoding method. The shape of gene feature:", train_gene_feature_onehotCoding.shape)
```

train\_gene\_feature\_onehotCoding is converted feature using one-hot encoding method. The shape of gene feature: (2124, 226)

#### Q4. How good is this gene feature in predicting $y_i$ ?

There are many ways to estimate how good a feature is, in predicting  $y_i$ . One of the good methods is to build a proper ML model using just this feature. In this case, we will build a logistic regression model using only Gene feature (one hot encoded) to predict  $y_i$ .

In [26]:

```
alpha = [10 ** x for x in range(-5, 1)] # hyperparam for SGD classifier.

# read more about SGDClassifier() at http://scikit-learn.org/stable/modules/generated/sklearn.linear_model.SGDClassifier.html
# -----
# default parameters
# SGDClassifier(loss='hinge', penalty='l2', alpha=0.0001, l1_ratio=0.15, fit_intercept=True, max_iter=None, tol=None,
# shuffle=True, verbose=0, epsilon=0.1, n_jobs=1, random_state=None, learning_rate='optimal', eta0=0.0, power_t=0.5,
# class_weight=None, warm_start=False, average=False, n_iter=None)

# some of methods
# fit(X, y[, coef_init, intercept_init, ...]) Fit linear model with Stochastic Gradient Descent.
# predict(X) Predict class labels for samples in X.

#-----
# video link:
#-----

cv_log_error_array=[]
for i in alpha:
    clf = SGDClassifier(alpha=i, penalty='l2', loss='log', random_state=42)
    clf.fit(train_gene_feature_onehotCoding, y_train)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_gene_feature_onehotCoding, y_train)
```

```

predict_y = sig_clf.predict_proba(cv_gene_feature_onehotCoding)
cv_log_error_array.append(log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))
print('For values of alpha = ', i, "The log loss is:", log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))


fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array, c='g')
for i, txt in enumerate(np.round(cv_log_error_array, 3)):
    ax.annotate((alpha[i], np.round(txt, 3)), (alpha[i], cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

best_alpha = np.argmin(cv_log_error_array)
clf = SGDCClassifier(alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=42)
clf.fit(train_gene_feature_onehotCoding, y_train)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_gene_feature_onehotCoding, y_train)

predict_y = sig_clf.predict_proba(train_gene_feature_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:", log_loss(y_train, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(cv_gene_feature_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The cross validation log loss is:", log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(test_gene_feature_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:", log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))

```

For values of alpha = 1e-05 The log loss is: 1.3368009935888565  
 For values of alpha = 0.0001 The log loss is: 1.1545038781790566  
 For values of alpha = 0.001 The log loss is: 1.1722262981090035  
 For values of alpha = 0.01 The log loss is: 1.306082679098516  
 For values of alpha = 0.1 The log loss is: 1.426237301488987  
 For values of alpha = 1 The log loss is: 1.4660850197925832



For values of best alpha = 0.0001 The train log loss is: 1.0387201537414488  
 For values of best alpha = 0.0001 The cross validation log loss is: 1.1545038781790566  
 For values of best alpha = 0.0001 The test log loss is: 1.2268569304490295

**Q5.** Is the Gene feature stable across all the data sets (Test, Train, Cross validation)?

**Ans.** Yes, it is. Otherwise, the CV and Test errors would be significantly more than train error.

In [27]:

```

print("Q6. How many data points in Test and CV datasets are covered by the ", unique_genes.shape[0], " genes in train dataset?")

test_coverage=test_df[test_df['Gene'].isin(list(set(train_df['Gene'])))].shape[0]
cv_coverage=cv_df[cv_df['Gene'].isin(list(set(train_df['Gene'])))].shape[0]

print('Ans\n1. In test data', test_coverage, 'out of', test_df.shape[0], ":", (test_coverage/test_df.

```

```
shape[0])*100)
print('2. In cross validation data',cv_coverage, 'out of ',cv_df.shape[0],":", (cv_coverage/cv_df.s
hape[0])*100)
```

Q6. How many data points in Test and CV datasets are covered by the 226 genes in train dataset?

Ans

1. In test data 635 out of 665 : 95.48872180451127
2. In cross validation data 515 out of 532 : 96.80451127819549

### 3.2.2 Univariate Analysis on Variation Feature

**Q7.** Variation, What type of feature is it ?

**Ans.** Variation is a categorical variable

**Q8.** How many categories are there?

In [28]:

```
unique_variations = train_df['Variation'].value_counts()
print('Number of Unique Variations :', unique_variations.shape[0])
# the top 10 variations that occurred most
print(unique_variations.head(10))
```

```
Number of Unique Variations : 1938
Truncating_Mutations      60
Deletion                   42
Amplification              42
Fusions                    21
G12V                       4
Q61H                       3
E17K                       3
G12D                       2
Overexpression             2
Q22K                       2
Name: Variation, dtype: int64
```


In [29]:


```
print("Ans: There are", unique_variations.shape[0], "different categories of variations in the
train data, and they are distributed as follows",)
```

Ans: There are 1938 different categories of variations in the train data, and they are distributed as follows

In [30]:

```
s = sum(unique_variations.values);
h = unique_variations.values/s;
plt.plot(h, label="Histogram of Variations")
plt.xlabel('Index of a Variation')
plt.ylabel('Number of Occurances')
plt.legend()
plt.grid()
plt.show()
```






In [31]:

```
c = np.cumsum(h)
print(c)
plt.plot(c, label='Cumulative distribution of Variations')
plt.grid()
plt.legend()
plt.show()
```

```
[0.02824859 0.0480226 0.06779661 ... 0.99905838 0.99952919 1. ]
```



### Q9. How to featurize this Variation feature ?

**Ans.** There are two ways we can featurize this variable check out this video:

<https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/handling-categorical-and-numerical-features/>

1. One hot Encoding
2. Response coding

We will be using both these methods to featurize the Variation Feature

In [32]:

```
# alpha is used for laplace smoothing
alpha = 1
# train gene feature
train_variation_feature_responseCoding = np.array(get_gv_feature(alpha, "Variation", train_df))
# test gene feature
test_variation_feature_responseCoding = np.array(get_gv_feature(alpha, "Variation", test_df))
# cross validation gene feature
cv_variation_feature_responseCoding = np.array(get_gv_feature(alpha, "Variation", cv_df))
```

In [33]:

```
print("train_variation_feature_responseCoding is a converted feature using the response coding method. The shape of Variation feature:", train_variation_feature_responseCoding.shape)
```

train\_variation\_feature\_responseCoding is a converted feature using the response coding method. The shape of Variation feature: (2124, 9)

In [34]:

```
# one-hot encoding of variation feature.
variation_vectorizer = TfidfVectorizer(max_features=2000, ngram_range=(1,4))
train_variation_feature_onehotCoding = variation_vectorizer.fit_transform(train_df['Variation'])
test_variation_feature_onehotCoding = variation_vectorizer.transform(test_df['Variation'])
cv_variation_feature_onehotCoding = variation_vectorizer.transform(cv_df['Variation'])
```

In [35]:

In [35]:

```
print("train_variation_feature_onehotEncoded is converted feature using the onne-hot encoding meth  
od. The shape of Variation feature:", train_variation_feature_onehotCoding.shape)
```

train\_variation\_feature\_onehotEncoded is converted feature using the onne-hot encoding method. The shape of Variation feature: (2124, 2000)

## Q10. How good is this Variation feature in predicting y\_i?


Let's build a model just like the earlier!

In [36]:

```
alpha = [10 ** x for x in range(-5, 1)]

# read more about SGDClassifier() at http://scikit-  
learn.org/stable/modules/generated/sklearn.linear_model.SGDClassifier.html  
# -----  
# default parameters  
# SGDClassifier(loss='hinge', penalty='l2', alpha=0.0001, l1_ratio=0.15, fit_intercept=True, max_i  
ter=None, tol=None,  
# shuffle=True, verbose=0, epsilon=0.1, n_jobs=1, random_state=None, learning_rate='optimal', eta0  
=0.0, power_t=0.5,  
# class_weight=None, warm_start=False, average=False, n_iter=None)  
  
# some of methods  
# fit(X, y[, coef_init, intercept_init, ...]) Fit linear model with Stochastic Gradient Descent.  
# predict(X) Predict class labels for samples in X.  
  
#-----  
# video link:  
#-----  
  
cv_log_error_array=[]  
for i in alpha:  
    clf = SGDClassifier(alpha=i, penalty='l2', loss='log', random_state=42)  
    clf.fit(train_variation_feature_onehotCoding, y_train)  
  
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")  
    sig_clf.fit(train_variation_feature_onehotCoding, y_train)  
    predict_y = sig_clf.predict_proba(cv_variation_feature_onehotCoding)  
  
    cv_log_error_array.append(log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))  
    print('For values of alpha = ', i, "The log loss is:", log_loss(y_cv, predict_y, labels=clf.clas  
ses_, eps=1e-15))  
  
fig, ax = plt.subplots()  
ax.plot(alpha, cv_log_error_array, c='g')  
for i, txt in enumerate(np.round(cv_log_error_array, 3)):  
    ax.annotate((alpha[i], np.round(txt, 3)), (alpha[i], cv_log_error_array[i]))  
plt.grid()  
plt.title("Cross Validation Error for each alpha")  
plt.xlabel("Alpha i's")  
plt.ylabel("Error measure")  
plt.show()  
  
best_alpha = np.argmin(cv_log_error_array)  
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=42)  
clf.fit(train_variation_feature_onehotCoding, y_train)  
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")  
sig_clf.fit(train_variation_feature_onehotCoding, y_train)  
  
predict_y = sig_clf.predict_proba(train_variation_feature_onehotCoding)  
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:", log_loss(y_train,  
predict_y, labels=clf.classes_, eps=1e-15))  
predict_y = sig_clf.predict_proba(cv_variation_feature_onehotCoding)  
print('For values of best alpha = ', alpha[best_alpha], "The cross validation log loss is:", log_lo  
ss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))  
predict_y = sig_clf.predict_proba(test_variation_feature_onehotCoding)  
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:", log_loss(y_test, p  
redict_y, labels=clf.classes_, eps=1e-15))
```

For values of alpha = 1e-05 The log loss is: 1.721544302989071  
 For values of alpha = 0.0001 The log loss is: 1.7092817793612023  
 For values of alpha = 0.001 The log loss is: 1.7193691732666683  
 For values of alpha = 0.01 The log loss is: 1.729206230678092  
 For values of alpha = 0.1 The log loss is: 1.741290479737912  
 For values of alpha = 1 The log loss is: 1.74208701536609



For values of best alpha = 0.0001 The train log loss is: 0.7627882666186675  
 For values of best alpha = 0.0001 The cross validation log loss is: 1.7092817793612023  
 For values of best alpha = 0.0001 The test log loss is: 1.684736874106243

**Q11. Is the Variation feature stable across all the data sets (Test, Train, Cross validation)?**

**Ans.** Not sure! But lets be very sure using the below analysis.

In [37]:

```
print("Q12. How many data points are covered by total ", unique_variations.shape[0], " genes in te
st and cross validation data sets?")
test_coverage=test_df[test_df['Variation'].isin(list(set(train_df['Variation'])))].shape[0]
cv_coverage=cv_df[cv_df['Variation'].isin(list(set(train_df['Variation'])))].shape[0]
print('Ans\n1. In test data',test_coverage, 'out of',test_df.shape[0], ":", (test_coverage/test_df.
shape[0])*100)
print('2. In cross validation data',cv_coverage, 'out of ',cv_df.shape[0]," :", (cv_coverage/cv_df.s
hape[0])*100)
```

Q12. How many data points are covered by total 1938 genes in test and cross validation data sets?

Ans

1. In test data 78 out of 665 : 11.729323308270677
2. In cross validation data 55 out of 532 : 10.338345864661653

### 3.2.3 Univariate Analysis on Text Feature

1. How many unique words are present in train data?
2. How are word frequencies distributed?
3. How to featurize text field?
4. Is the text feature useful in predicting  $y_i$ ?
5. Is the text feature stable across train, test and CV datasets?

In [38]:

```
# cls_text is a data frame
# for every row in data fram consider the 'TEXT'
# split the words by space
# make a dict with those words
# increment its count whenever we see that word

def extract_dictionary_paddle(cls_text):
    dictionary = defaultdict(int)
    for index, row in cls_text.iterrows():
        for word in row['TEXT'].split():
```

```

        dictionary[word] +=1
    return dictionary

```

In [39]:

```

import math
#https://stackoverflow.com/a/1602964
def get_text_responsecoding(df):
    text_feature_responseCoding = np.zeros((df.shape[0],9))
    for i in range(0,9):
        row_index = 0
        for index, row in df.iterrows():
            sum_prob = 0
            for word in row['TEXT'].split():
                sum_prob += math.log(((dict_list[i].get(word,0)+10 )/(total_dict.get(word,0)+90)))
            text_feature_responseCoding[row_index][i] = math.exp(sum_prob/len(row['TEXT'].split()))
            row_index += 1
    return text_feature_responseCoding

```

In [40]:

```

# building a TfidfVectorizer with all the words that occurred minimum 3 times in train data
#text_vectorizer = CountVectorizer(min_df=3)
text_vectorizer = TfidfVectorizer(min_df=3, max_features=2000, ngram_range=(1,4))
train_text_feature_onehotCoding = text_vectorizer.fit_transform(train_df['TEXT'])
# getting all the feature names (words)
train_text_features= text_vectorizer.get_feature_names()

# train_text_feature_onehotCoding.sum(axis=0).A1 will sum every row and returns (1*number of features) vector
train_text_fea_counts = train_text_feature_onehotCoding.sum(axis=0).A1

# zip(list(text_features),text_fea_counts) will zip a word with its number of times it occurred
text_fea_dict = dict(zip(list(train_text_features),train_text_fea_counts))

print("Total number of unique words in train data :", len(train_text_features))

```

Total number of unique words in train data : 2000

In [41]:

```

dict_list = []
# dict_list =[] contains 9 dictionaries each corresponds to a class
for i in range(1,10):
    cls_text = train_df[train_df['Class']==i]
    # build a word dict based on the words in that class
    dict_list.append(extract_dictionary_paddle(cls_text))
    # append it to dict_list

# dict_list[i] is build on i'th class text data
# total_dict is build on whole training text data
total_dict = extract_dictionary_paddle(train_df)

confuse_array = []
for i in train_text_features:
    ratios = []
    max_val = -1
    for j in range(0,9):
        ratios.append((dict_list[j][i]+10 )/(total_dict[i]+90))
    confuse_array.append(ratios)
confuse_array = np.array(confuse_array)

```

In [42]:

```

#response coding of text features
train_text_feature_responseCoding = get_text_responsecoding(train_df)
test_text_feature_responseCoding = get_text_responsecoding(test_df)
cv_text_feature_responseCoding = get_text_responsecoding(cv_df)

```

In [43]:

```
# https://stackoverflow.com/a/16202486
# we convert each row values such that they sum to 1
train_text_feature_responseCoding =
(train_text_feature_responseCoding.T/train_text_feature_responseCoding.sum(axis=1)).T
test_text_feature_responseCoding =
(test_text_feature_responseCoding.T/test_text_feature_responseCoding.sum(axis=1)).T
cv_text_feature_responseCoding = (cv_text_feature_responseCoding.T/cv_text_feature_responseCoding.
sum(axis=1)).T
```

In [44]:

```
# don't forget to normalize every feature
train_text_feature_onehotCoding = normalize(train_text_feature_onehotCoding, axis=0)

# we use the same vectorizer that was trained on train data
test_text_feature_onehotCoding = text_vectorizer.transform(test_df['TEXT'])
# don't forget to normalize every feature
test_text_feature_onehotCoding = normalize(test_text_feature_onehotCoding, axis=0)

# we use the same vectorizer that was trained on train data
cv_text_feature_onehotCoding = text_vectorizer.transform(cv_df['TEXT'])
# don't forget to normalize every feature
cv_text_feature_onehotCoding = normalize(cv_text_feature_onehotCoding, axis=0)
```

In [45]:

```
#https://stackoverflow.com/a/2258273/4084039
sorted_text_fea_dict = dict(sorted(text_fea_dict.items(), key=lambda x: x[1] , reverse=True))
sorted_text_occur = np.array(list(sorted_text_fea_dict.values()))
```

In [46]:

```
# Number of words for a given frequency.
print(Counter(sorted_text_occur))
```

```
Counter({7.555398648475674: 6, 11.333097972713519: 5, 6.084890017977853: 3, 5.837514862114463: 3,
9.569432890443561: 2, 9.518362723771245: 2, 8.662758983111265: 2, 8.609353712515176: 2,
7.880821430931938: 2, 7.73899525742619: 2, 7.66302878608496: 2, 7.480755829437548: 2,
7.252894333930967: 2, 7.191947043264201: 2, 6.79209989045965: 2, 6.116003259966491: 2,
6.11510297614852: 2, 210.17213732995248: 1, 142.3344221340105: 1, 123.28194520145317: 1,
110.30418050301618: 1, 101.54465790026588: 1, 96.40472086250529: 1, 92.21742622841388: 1,
92.16101388282378: 1, 92.03230838856803: 1, 90.21216689199295: 1, 89.17273757478354: 1,
86.51914003681736: 1, 76.04868758125683: 1, 75.32989465712537: 1, 72.6891258269478: 1,
71.8719932797566: 1, 71.12333207770874: 1, 67.0227817177874: 1, 66.88678913705756: 1,
66.54969328051459: 1, 64.48349993936363: 1, 60.86916536471965: 1, 60.86269182118953: 1,
58.91526803327548: 1, 55.434265612555514: 1, 55.38085348052518: 1, 55.085635888586935: 1,
54.864714805269415: 1, 54.412964032500916: 1, 54.10138723992345: 1, 52.74932959448596: 1,
52.6541386385727: 1, 52.11630937364783: 1, 51.44174397437812: 1, 50.577206151704075: 1,
50.145678937930555: 1, 48.57102608341427: 1, 48.39632897412581: 1, 48.25993260883129: 1, 46.6610066
2205091: 1, 46.630087434876785: 1, 45.547233156586344: 1, 43.63972195595017: 1, 43.0267005459537:
1, 42.93001682596737: 1, 42.48420782691775: 1, 42.47832303069629: 1, 40.535716417935326: 1,
40.09598343712877: 1, 38.169346512934155: 1, 38.06332923625239: 1, 37.89501471450194: 1,
36.95358202433027: 1, 36.868128038705: 1, 36.05432343864839: 1, 35.92362186520374: 1,
35.80361040444916: 1, 35.77279697917545: 1, 35.726819578283006: 1, 35.275919953665024: 1,
35.22642749317283: 1, 35.01720476013494: 1, 34.97841475996995: 1, 34.58995809399706: 1,
33.934037312182575: 1, 33.730051233951045: 1, 33.406731749861294: 1, 33.37919897051788: 1,
33.19317152647789: 1, 32.947712776199644: 1, 32.62121471096092: 1, 32.261802643365385: 1,
32.10991962230041: 1, 31.872832584532034: 1, 31.72275219685889: 1, 31.679270026395283: 1,
31.650725552911716: 1, 30.896230907188034: 1, 30.646685546054524: 1, 30.35617836012703: 1,
30.29163030353833: 1, 30.193566186535143: 1, 30.064679387105173: 1, 29.826751083953: 1,
29.512297400093487: 1, 29.174635564854523: 1, 29.113419298377686: 1, 28.977463141272924: 1,
28.96931889066878: 1, 28.8304198747675: 1, 28.532549374067774: 1, 28.416639514112486: 1,
28.386095946085465: 1, 27.981025316847347: 1, 27.862124712200643: 1, 27.409867054053866: 1,
27.196067077821343: 1, 27.03864964463922: 1, 26.882475643073573: 1, 26.869072314199947: 1,
26.73015517440322: 1, 26.714079574551118: 1, 26.710391977805692: 1, 26.7030509876286: 1,
26.7001985752347: 1, 26.492866254035725: 1, 26.38533414989864: 1, 25.86743099225329: 1,
25.857816545998904: 1, 25.75109838983151: 1, 25.657851403050667: 1, 25.62230133307364: 1,
25.616105781032218: 1, 25.582461168234776: 1, 25.527870217651163: 1, 25.366949404652615: 1,
25.296965747683885: 1, 25.16922703830987: 1, 25.16233006055522: 1, 25.15969596638238: 1, 25.0834908
01690353: 1, 25.010422340292667: 1, 24.992477816297615: 1, 24.929286127284794: 1,
24.852814298483324: 1, 24.516222731925385: 1, 24.181571418749353: 1, 24.126987622585926: 1,
24.099407660464703: 1, 24.094469332927893: 1, 24.065940900282325: 1, 24.020536646628237: 1,
```



21.055407000004705: 1, 21.05540700000221055: 1, 21.0554070000020220: 1, 21.0205540700020257: 1, 23.81979167063137: 1, 23.714049759626214: 1, 23.69942516439738: 1, 23.687172531307368: 1, 23.68053850391787: 1, 23.64507108254561: 1, 23.59322494161485: 1, 23.4864246940459: 1, 23.425743010224167: 1, 23.389623763259827: 1, 23.262531804979332: 1, 23.085370559519223: 1, 22.91339829763365: 1, 22.688764186937558: 1, 22.652070391529946: 1, 22.61730783452018: 1, 22.532825854500693: 1, 22.397075486623173: 1, 22.331073395580848: 1, 22.240811915244098: 1, 22.03207730507138: 1, 21.771504228946426: 1, 21.760040909825975: 1, 21.740306855411415: 1, 21.608111482169583: 1, 21.506293124299006: 1, 21.338773810867888: 1, 21.208354351088904: 1, 20.84567007221386: 1, 20.820285228166526: 1, 20.626652042136918: 1, 20.567710206744042: 1, 20.430274501990286: 1, 20.387983041688514: 1, 20.35437436532027: 1, 20.348058618720632: 1, 20.244900964968185: 1, 20.20886894886593: 1, 20.18681196179705: 1, 20.162900140643856: 1, 20.121330489287754: 1, 20.06647478523907: 1, 20.02843885673455: 1, 20.01962137772522: 1, 20.0013574 81479157: 1, 19.986863803972927: 1, 19.964605671963046: 1, 19.952307680942454: 1, 19.787200889589716: 1, 19.730349675744513: 1, 19.618131920494307: 1, 19.56024394110674: 1, 19.437454097295202: 1, 19.33783232941244: 1, 19.33622019353207: 1, 19.299544459210566: 1, 19.248616399631796: 1, 19.226509717112087: 1, 19.202623903587554: 1, 19.150731583808962: 1, 19.10084687796415: 1, 19.069141602965107: 1, 18.936385856302483: 1, 18.836320414373844: 1, 18.82003388919716: 1, 18.80878242119015: 1, 18.639288718513654: 1, 18.62669405406039: 1, 18.614632353708423: 1, 18.606441081625018: 1, 18.588953299794557: 1, 18.579341129393143: 1, 18.46080568246849: 1, 18.452934628507567: 1, 18.42518775540271: 1, 18.365120239933248: 1, 18.356753663344108: 1, 18.278199107472915: 1, 18.27196817100119: 1, 18.252146241489875: 1, 18.174054848791265: 1, 18.166838766170777: 1, 18.133331727200364: 1, 18.08868322279582: 1, 18.048934521197957: 1, 18.040734110279242: 1, 17.988250428232643: 1, 17.97188153360032: 1, 17.953301265637574: 1, 17.933415320272125: 1, 17.87730438342532: 1, 17.835223177244274: 1, 17.82428202555092: 1, 17.81905404786737: 1, 17.799461656789113: 1, 17.787219015955984: 1, 17.786999205472554: 1, 17.775792015737505: 1, 17.733212156517233: 1, 17.73217514917421: 1, 17.72455460961499: 1, 17.66836950200135: 1, 17.650107332931842: 1, 17.615486072888242: 1, 17.5705380269576: 1, 17.51217844502385: 1, 17.440161677854924: 1, 17.439332625685996: 1, 17.42962218156537: 1, 17.398003948786165: 1, 17.337723873414692: 1, 17.327837876957954: 1, 17.32553217674111: 1, 17.321093856151663: 1, 17.31250724275993: 1, 17.2512950763265: 1, 17.229092743305706: 1, 17.132790700770798: 1, 17.093640342634274: 1, 17.070909057106974: 1, 17.068986260333734: 1, 16.949544540048688: 1, 16.925639380011944: 1, 16.832333181501745: 1, 16.799700821855: 1, 16.74485253000229: 1, 16.72978972456953: 1, 16.64546462063524: 1, 16.6441321418 98684: 1, 16.60705020835447: 1, 16.594575264880955: 1, 16.55512562763352: 1, 16.505436747294514: 1, 16.449628606681063: 1, 16.428992034644295: 1, 16.428832282818295: 1, 16.36369880799406: 1, 16.325411384697805: 1, 16.28056064003108: 1, 16.275155004185496: 1, 16.219771154578716: 1, 16.176775279630318: 1, 16.12416334548479: 1, 16.09834443134257: 1, 16.071376162624915: 1, 16.070516202401176: 1, 16.01348519030774: 1, 15.991129047404751: 1, 15.990216298178439: 1, 15.978141255757127: 1, 15.90111342394248: 1, 15.899691398663174: 1, 15.855060810887704: 1, 15.841973490766565: 1, 15.767044610252233: 1, 15.765370511542384: 1, 15.752051056104953: 1, 15.745332032367436: 1, 15.726978980064903: 1, 15.722755975384079: 1, 15.712055956330214: 1, 15.70521572548607: 1, 15.67515806673969: 1, 15.591507467360842: 1, 15.555454356766203: 1, 15.52850351151423: 1, 15.476338976540994: 1, 15.451998606194826: 1, 15.44921790234777: 1, 15.442373271838354: 1, 15.399186692077091: 1, 15.387056474004794: 1, 15.345015742411793: 1, 15.275198428939513: 1, 15.25485069765909: 1, 15.241692447849557: 1, 15.219011957496626: 1, 15.100339840453293: 1, 15.064392273312409: 1, 15.058594204731445: 1, 15.02171216091215: 1, 15.018370831754268: 1, 15.01050458129975: 1, 14.977348272278054: 1, 14.944001309315633: 1, 14.918398737739786: 1, 14.895408410417014: 1, 14.894395347197625: 1, 14.794500380530847: 1, 14.792914689655849: 1, 14.757906570189796: 1, 14.756712456596386: 1, 14.713681522891441: 1, 14.707542648888566: 1, 14.701119769065347: 1, 14.661279708515435: 1, 14.638353845064442: 1, 14.617252256409477: 1, 14.612569480776282: 1, 14.600773718357424: 1, 14.598286406673635: 1, 14.527701254322622: 1, 14.503071673474844: 1, 14.476652408553083: 1, 14.46424988720096: 1, 14.4347640261633: 1, 14.406731353700017: 1, 14.386541264041572: 1, 14.37982224792086: 1, 14.363556227023608: 1, 14.315929241141369: 1, 14.31493515981519: 1, 14.31206008382544: 1, 14.299128133601386: 1, 14.219760787446367: 1, 14.200546406260388: 1, 14.175951739646546: 1, 14.119869161436343: 1, 14.077189769726193: 1, 14.049896398984766: 1, 14.049199754059691: 1, 14.042629312045205: 1, 14.009786859472351: 1, 13.98419614479772: 1, 13.933384397343406: 1, 13.927474327836212: 1, 13.906540555342092: 1, 13.89534628383061: 1, 13.8586085362376: 1, 13.851556400391965: 1, 13.851025008933116: 1, 13.83358108101166: 1, 13.829492109698489: 1, 13.824154103985167: 1, 13.781280887419683: 1, 13.75916488330976: 1, 13.753766172248488: 1, 13.751965225862056: 1, 13.750365479276699: 1, 13.74007818306975: 1, 13.6799817488826: 1, 13.679680740921528: 1, 13.675998227498503: 1, 13.642945726758434: 1, 13.632091913953358: 1, 13.602177020155517: 1, 13.56267354604367: 1, 13.559065679929247: 1, 13.520868613536045: 1, 13.483704360287982: 1, 13.477226886166262: 1, 13.472138923151395: 1, 13.454100368859983: 1, 13.433984771858198: 1, 13.41944675751567: 1, 13.39772796278204: 1, 13.35891553363707: 1, 13.3184942284506: 1, 13.312864511290611: 1, 13.271838490711907: 1, 13.24544867481554: 1, 13.242121754713759: 1, 13.23317475655189: 1, 13.207825430760979: 1, 13.177030806961987: 1, 13.156963235485602: 1, 13.133079774275789: 1, 13.126877876994731: 1, 13.125711877061097: 1, 13.095015704578211: 1, 13.082651470497538: 1, 13.07809687088851: 1, 13.040558921320413: 1, 12.976954553439347: 1, 12.962666738716216: 1, 12.949208310863378: 1, 12.932759069894882: 1, 12.89235557847474: 1, 12.877005083082041: 1, 12.873478220952332: 1, 12.872460772390099: 1, 12.869047140054102: 1, 12.860684052367956: 1, 12.835606965904567: 1, 12.812786756090121: 1, 12.767909928567288: 1, 12.759928255642315: 1, 12.739755019989007: 1, 12.717007009434328: 1, 12.681633613914927: 1, 12.677881637075906: 1, 12.665176724624631: 1, 12.631269748828936: 1, 12.63057541272834: 1, 12.616950254773416: 1, 12.611264082996167: 1, 12.609294510091084: 1, 12.604750928596506: 1, 12.579319637349334: 1, 12.568292318400989: 1, 12.526228533428828: 1, 12.482475935047448: 1, 12.475788331429856: 1, 12.44114889219502: 1, 12.437163892131121: 1, 12.410253410705712: 1, 12.406832210318658: 1, 12.3827951963220368: 1, 12.383893674930059: 1

12.410233410705712: 1, 12.400032210310030: 1, 12.392793170320300: 1, 12.303039074930033: 1, 12.371216277641736: 1, 12.361333160847735: 1, 12.360487836646458: 1, 12.357041125417265: 1, 12.33831197204751: 1, 12.333623582858536: 1, 12.30514032140291: 1, 12.291516507808549: 1, 12.290576801688742: 1, 12.257372992647642: 1, 12.249336504912348: 1, 12.24776322678506: 1, 12.242429840829798: 1, 12.228005332871065: 1, 12.2258987876097: 1, 12.205507536709284: 1, 12.197826072555308: 1, 12.189371835352963: 1, 12.170339491396422: 1, 12.1221688614411: 1, 12.095233254212173: 1, 12.050728447655864: 1, 12.044120299562858: 1, 12.018294466810868: 1, 12.005124920054271: 1, 11.978314053408328: 1, 11.964035178097252: 1, 11.960946351979896: 1, 11.958586665009973: 1, 11.94587173066547: 1, 11.943195741162068: 1, 11.93155377262226: 1, 11.927682278872341: 1, 11.902216089563833: 1, 11.902161491669142: 1, 11.901976506515254: 1, 11.886640935861635: 1, 11.864031883950931: 1, 11.848488784058997: 1, 11.840936956395947: 1, 11.823275355619858: 1, 11.812661380727214: 1, 11.794198967048336: 1, 11.7915996860871: 1, 11.785357490918381: 1, 11.767499111103964: 1, 11.766011703245994: 1, 11.759163163282379: 1, 11.753296316955273: 1, 11.751198095172718: 1, 11.736017940624382: 1, 11.734679348881421: 1, 11.731771158127923: 1, 11.724211382020645: 1, 11.712378927332447: 1, 11.704742060098267: 1, 11.701278359325299: 1, 11.668769444655712: 1, 11.667359425272418: 1, 11.65309864314746: 1, 11.64564320638673: 1, 11.627105186271939: 1, 11.605056484028822: 1, 11.59583556762811: 1, 11.58083572108677: 1, 11.574828054895912: 1, 11.550073239108317: 1, 11.5423360347067765: 1, 11.514968418168568: 1, 11.509807992950272: 1, 11.475625859285534: 1, 11.47255442608459: 1, 11.438604878480412: 1, 11.435766838219948: 1, 11.411205635677058: 1, 11.367989388491477: 1, 11.367434386811441: 1, 11.365086334812768: 1, 11.33283045548382: 1, 11.32407576623401: 1, 11.321242180262413: 1, 11.306436851778962: 1, 11.26506385299157: 1, 11.197434355039894: 1, 11.194105424147494: 1, 11.192523198821874: 1, 11.17593457592564: 1, 11.175400397591197: 1, 11.125538652242245: 1, 11.116699356223174: 1, 11.112274021963092: 1, 11.104055565378014: 1, 11.08551664566608: 1, 11.085061131420893: 1, 11.066725382332692: 1, 11.056791769508886: 1, 11.002127589647936: 1, 10.998679896329778: 1, 10.991026688247615: 1, 10.986986280638824: 1, 10.985447373983405: 1, 10.961484988988406: 1, 10.95809597430431: 1, 10.946480687136884: 1, 10.942639599204643: 1, 10.892368768021495: 1, 10.885414542780513: 1, 10.87817143684149: 1, 10.86140443947583: 1, 10.8488760824418: 1, 10.801598821996924: 1, 10.788822281708581: 1, 10.784041862347715: 1, 10.776237589721077: 1, 10.753691233504762: 1, 10.743570046602045: 1, 10.74049377874696: 1, 10.738220911189368: 1, 10.733477813082462: 1, 10.723713915884865: 1, 10.718933900808784: 1, 10.717289452270375: 1, 10.714249089767304: 1, 10.710885070708065: 1, 10.706068687335971: 1, 10.691926441309816: 1, 10.680227840335354: 1, 10.677489976261812: 1, 10.637586252666752: 1, 10.635376304864156: 1, 10.59051133340477: 1, 10.566642945501938: 1, 10.551818360506195: 1, 10.54558995623577: 1, 10.54379247040237: 1, 10.543750891224864: 1, 10.538706349391724: 1, 10.529513106872846: 1, 10.52128029000255: 1, 10.513583297831664: 1, 10.51308633081296: 1, 10.509430699576265: 1, 10.473037342609326: 1, 10.46472320888208: 1, 10.430890275158255: 1, 10.426470663047125: 1, 10.404634644520483: 1, 10.374684573190184: 1, 10.371306581097947: 1, 10.37081113925501: 1, 10.364034083092347: 1, 10.33417521509407: 1, 10.323453552416748: 1, 10.31844888383058: 1, 10.30270994304847: 1, 10.285523931283223: 1, 10.283363422960237: 1, 10.211905157571426: 1, 10.20233406917091: 1, 10.198062272206801: 1, 10.189294898621805: 1, 10.181839385955268: 1, 10.165976197786586: 1, 10.152555913413368: 1, 10.107779916668932: 1, 10.093438011193907: 1, 10.040474445489838: 1, 10.038621948020447: 1, 10.01721620191406: 1, 10.011413734880895: 1, 10.006215727527444: 1, 9.995678311171886: 1, 9.985713257002578: 1, 9.981694140979302: 1, 9.978363254971049: 1, 9.960857520596036: 1, 9.95742149431632: 1, 9.9538680212865: 1, 9.93918270552297: 1, 9.93761957067938: 1, 9.91347566700546: 1, 9.906697452008249: 1, 9.906607506481969: 1, 9.902031896141567: 1, 9.8913760366657: 1, 9.879824921628009: 1, 9.878730947814974: 1, 9.869648419107884: 1, 9.865006397651255: 1, 9.851538843430534: 1, 9.84446604513455: 1, 9.835464455250971: 1, 9.833762092459155: 1, 9.824678186980153: 1, 9.82162475179499: 1, 9.810248654325935: 1, 9.795555079907611: 1, 9.770494933316316: 1, 9.755041077631818: 1, 9.735686975580508: 1, 9.734709383907514: 1, 9.673147298397321: 1, 9.668049122981527: 1, 9.665606212664633: 1, 9.646689649099908: 1, 9.634415664855894: 1, 9.609855473426654: 1, 9.59655609105959: 1, 9.59603571249139: 1, 9.593775662187603: 1, 9.591159718374193: 1, 9.580991243633031: 1, 9.566933005400927: 1, 9.56112324081851: 1, 9.558577351943505: 1, 9.552403525012034: 1, 9.547072451267548: 1, 9.527473949496528: 1, 9.52363253187674: 1, 9.519405985709197: 1, 9.51398151961645: 1, 9.481007817083945: 1, 9.47215247494477: 1, 9.46947998356316: 1, 9.45938555901646: 1, 9.454707759406693: 1, 9.449188919583488: 1, 9.413109276905066: 1, 9.409327218224979: 1, 9.405701986736211: 1, 9.403157235127718: 1, 9.401871522586566: 1, 9.40117645958275: 1, 9.399494889670072: 1, 9.39848191928965: 1, 9.39769749117273: 1, 9.3932485330912: 1, 9.390302014655001: 1, 9.383006902006771: 1, 9.369171976280935: 1, 9.365606522609534: 1, 9.35081752725456: 1, 9.339047849472491: 1, 9.332083568368876: 1, 9.331283101633039: 1, 9.303119614158824: 1, 9.294087160447685: 1, 9.283712725531236: 1, 9.263983642711201: 1, 9.263779499475707: 1, 9.252532068880026: 1, 9.241151390590396: 1, 9.240510451108719: 1, 9.224165177886453: 1, 9.223719164362372: 1, 9.215934238515553: 1, 9.209936778513589: 1, 9.178033627567709: 1, 9.166676322913567: 1, 9.15733726681232: 1, 9.156759578852022: 1, 9.155104210270707: 1, 9.150313410086435: 1, 9.148783767311766: 1, 9.14323430608329: 1, 9.112405560319234: 1, 9.111853166944417: 1, 9.110869637074265: 1, 9.09627374992923: 1, 9.086238669829724: 1, 9.074120357889019: 1, 9.071417708315671: 1, 9.07066860218776: 1, 9.052340488263205: 1, 9.04652895700805: 1, 9.04542165043859: 1, 9.044977379510971: 1, 9.039054560271348: 1, 9.038805498510417: 1, 9.035676498338129: 1, 9.022621875635162: 1, 9.017127207421032: 1, 9.01348435860604: 1, 9.00961076561612: 1, 9.004233409334065: 1, 9.00069233143473: 1, 8.987963714618415: 1, 8.979502926285548: 1, 8.964096946618998: 1, 8.962940697909245: 1, 8.95402015339493: 1, 8.948794581188254: 1, 8.948550500009114: 1, 8.944210610177894: 1, 8.940836701423756: 1, 8.93196351570033: 1, 8.925196524906866: 1, 8.92072259210571: 1, 8.916778846039481: 1, 8.91252003732672: 1, 8.907921837180274: 1, 8.90287910316152: 1, 8.900146433873775: 1, 8.900057517910833: 1, 8.898874839461614: 1, 8.895228114286700: 1, 8.890011921812006: 1, 8.878478752157054: 1, 8.875000472802104: 1

0.093530114200709: 1, 0.0900110219153990: 1, 0.07947052157034: 1, 0.073909472005194: 1,  
8.874976975203836: 1, 8.868242818076114: 1, 8.867310140721733: 1, 8.861670491585516: 1,  
8.856752158809247: 1, 8.851197353027324: 1, 8.845469897233466: 1, 8.838429433193172: 1,  
8.825232236084277: 1, 8.82249686142137: 1, 8.814064201716922: 1, 8.813402064710214: 1,  
8.798650891257388: 1, 8.791963897316426: 1, 8.779785046444864: 1, 8.761428187994085: 1,  
8.754377273039989: 1, 8.752299303572421: 1, 8.723947300389932: 1, 8.722006404969694: 1,  
8.719199756718238: 1, 8.713969138101021: 1, 8.692757786456626: 1, 8.689413422307918: 1,  
8.68879672647158: 1, 8.685076223636859: 1, 8.67679373855848: 1, 8.673010392484025: 1,  
8.67282121132028: 1, 8.666529786806635: 1, 8.653905346472877: 1, 8.647943398760752: 1,  
8.631684726078309: 1, 8.623796981531433: 1, 8.61449721827668: 1, 8.606609660619291: 1,  
8.589906744853385: 1, 8.52191257655843: 1, 8.52072226328442: 1, 8.520404963200093: 1,  
8.507533202857765: 1, 8.505415797791374: 1, 8.502603692853201: 1, 8.499538093967601: 1,  
8.493500175711965: 1, 8.485206066372566: 1, 8.475473073651571: 1, 8.475160180395179: 1,  
8.471767382858195: 1, 8.471471796592972: 1, 8.464584844157626: 1, 8.4478946908168: 1,  
8.446892020069706: 1, 8.42950449232296: 1, 8.429275463980446: 1, 8.428392048518774: 1,  
8.419288332930888: 1, 8.418426470385597: 1, 8.403969497288523: 1, 8.386940884991818: 1,  
8.380850237211995: 1, 8.378699571122468: 1, 8.375122056475002: 1, 8.366670896962221: 1,  
8.341606229386658: 1, 8.328110497721822: 1, 8.31729446742681: 1, 8.3172611478442: 1,  
8.314121427494394: 1, 8.298629334571563: 1, 8.296960154900793: 1, 8.288949210183976: 1,  
8.282127755995878: 1, 8.272006167066282: 1, 8.269085015563217: 1, 8.255988336112047: 1,  
8.248691903264413: 1, 8.239582260511307: 1, 8.23865565601418: 1, 8.23656948913147: 1,  
8.230632312582289: 1, 8.222446797180785: 1, 8.21870574655496: 1, 8.215109903547027: 1,  
8.212124968650652: 1, 8.19969655286386: 1, 8.198303892387944: 1, 8.191893703433017: 1,  
8.187454380164375: 1, 8.182494851526886: 1, 8.179095622293957: 1, 8.17432939614024: 1,  
8.172239954236984: 1, 8.169546926280326: 1, 8.144641833409775: 1, 8.14108089987498: 1,  
8.133097333036748: 1, 8.131337090457635: 1, 8.116790351372027: 1, 8.11365388255101: 1,  
8.106591300205222: 1, 8.103756190915773: 1, 8.102016184695383: 1, 8.090977311179554: 1,  
8.089703941524737: 1, 8.082377961905431: 1, 8.075894973051826: 1, 8.072224513236398: 1,  
8.06367310145897: 1, 8.048095568393395: 1, 8.043832439774652: 1, 8.038985116395462: 1,  
8.037802047135582: 1, 8.030622919017599: 1, 8.023270443934715: 1, 8.01507857504261: 1,  
8.013316487684195: 1, 8.01160652743312: 1, 8.01026987320927: 1, 8.005153271687679: 1,  
7.993159249185121: 1, 7.983771361264074: 1, 7.982543808410491: 1, 7.968708509820759: 1,  
7.933276673834268: 1, 7.932884759285903: 1, 7.909398088742683: 1, 7.896888333098004: 1,  
7.895426296920024: 1, 7.883081808170706: 1, 7.882552470597211: 1, 7.876312931781857: 1,  
7.873197907113511: 1, 7.867172724460702: 1, 7.861683032211123: 1, 7.858704934186268: 1,  
7.851400169237529: 1, 7.848616427812987: 1, 7.846224310924183: 1, 7.840885261680764: 1,  
7.835368932972791: 1, 7.831661430816821: 1, 7.830146271404744: 1, 7.827363161858427: 1,  
7.818024550581674: 1, 7.817261469918491: 1, 7.812861417449877: 1, 7.810528987449796: 1,  
7.809769332460079: 1, 7.780834142668084: 1, 7.780183838411205: 1, 7.777748128129155: 1,  
7.776934009898788: 1, 7.76814925140587: 1, 7.761695400628886: 1, 7.746070432463337: 1,  
7.7415129544588535: 1, 7.741263576904253: 1, 7.738548343087118: 1, 7.737966170697089: 1, 7.73459820  
3356802: 1, 7.729485507359188: 1, 7.7275165280899465: 1, 7.726540941514875: 1, 7.726171975805724:  
1, 7.725093211828159: 1, 7.723382408914496: 1, 7.722522397049822: 1, 7.722119556387749: 1,  
7.720972442240229: 1, 7.718985902434579: 1, 7.717577781891459: 1, 7.713267909609716: 1,  
7.712007689863777: 1, 7.708515187677371: 1, 7.705728302501317: 1, 7.6968202477966585: 1,  
7.695325451127633: 1, 7.692851677067502: 1, 7.690829937702188: 1, 7.686601749873006: 1,  
7.68637811959028: 1, 7.685019967604363: 1, 7.660267068405359: 1, 7.651717735835529: 1,  
7.641890151954027: 1, 7.640801913570627: 1, 7.623765870722352: 1, 7.596246905716373: 1,  
7.595006834368883: 1, 7.594922746392977: 1, 7.5933744547407045: 1, 7.5841784603785385: 1,  
7.58163432567336: 1, 7.580426381355497: 1, 7.5732606539358285: 1, 7.570206136763692: 1,  
7.5537810423168015: 1, 7.5448028300635785: 1, 7.5430429134257295: 1, 7.542232287332838: 1,  
7.534328464245136: 1, 7.518409708643686: 1, 7.517174798195781: 1, 7.504102207804057: 1,  
7.502947218302731: 1, 7.501484961979408: 1, 7.5012144733048425: 1, 7.495011653783422: 1,  
7.487320559547685: 1, 7.487235052763313: 1, 7.482725448371162: 1, 7.480342386743174: 1,  
7.479467957836917: 1, 7.478149555472529: 1, 7.47400673032077: 1, 7.469192941580896: 1,  
7.462033066447782: 1, 7.452921029733177: 1, 7.447923887258356: 1, 7.438530573040895: 1,  
7.430313834773704: 1, 7.425260446034159: 1, 7.414013464159051: 1, 7.409807826994777: 1,  
7.408809893443133: 1, 7.405083656527651: 1, 7.404309233482531: 1, 7.400987016397494: 1,  
7.3980956388450965: 1, 7.388699328411856: 1, 7.386080501914618: 1, 7.37438461725188: 1,  
7.372142381598839: 1, 7.3696143341678235: 1, 7.363728827795407: 1, 7.363211145670119: 1,  
7.36174192631139: 1, 7.361520627040107: 1, 7.3561712800579135: 1, 7.3530652552025115: 1,  
7.330781100990578: 1, 7.329438588145448: 1, 7.329079739255432: 1, 7.326433101454144: 1,  
7.320641274332687: 1, 7.317227254867722: 1, 7.312710462229368: 1, 7.301974375280389: 1,  
7.300272435943351: 1, 7.2977575195948425: 1, 7.292386543918528: 1, 7.290447872127987: 1,  
7.282374748278682: 1, 7.277217935818352: 1, 7.267311933181212: 1, 7.265132085960661: 1,  
7.2550541522321925: 1, 7.253699317880438: 1, 7.236885695069517: 1, 7.230907864221126: 1, 7.22880673  
7198312: 1, 7.224902298671739: 1, 7.224669998156096: 1, 7.223704722936156: 1, 7.219241968142705: 1,  
7.216315161358054: 1, 7.212936023955751: 1, 7.21026380040515: 1, 7.201305712309626: 1,  
7.1991686329015705: 1, 7.195763661194484: 1, 7.192594605362417: 1, 7.191074325368699: 1, 7.18894843  
0770205: 1, 7.181816709094684: 1, 7.180140106499083: 1, 7.178936822378236: 1, 7.1777925212581755:  
1, 7.167818774040833: 1, 7.16385998545013: 1, 7.148667018302534: 1, 7.147210584143498: 1,  
7.146300495640063: 1, 7.140915878718705: 1, 7.138043216945064: 1, 7.135469975869714: 1,  
7.125030864417373: 1, 7.11800361031413: 1, 7.115457985478031: 1, 7.113892746647061: 1,  
7.112379960857452: 1, 7.110504804506296: 1, 7.1092282039145145: 1, 7.107801614776538: 1,  
7.098817003870962: 1, 7.096062927162265: 1, 7.091547789778205: 1, 7.087451156719485: 1,  
7.08589208032998: 1, 7.082555845511112: 1, 7.075522126261094: 1, 7.074089960847658: 1,  
7.070351793404775: 1, 7.069691301567083: 1, 7.066354456824233: 1, 7.064114001144847: 1,  
7.062440251340772: 1, 7.0547323081610425: 1, 7.051822323042316: 1, 7.046200231417341: 1

7.063402351340772: 1, 7.0547271091610435: 1, 7.051830333943116: 1, 7.046392731417341: 1, 7.033590503193685: 1, 7.022276188684467: 1, 7.006107054614596: 1, 6.996749909800846: 1, 6.995660916112889: 1, 6.989796451846528: 1, 6.981138081128712: 1, 6.9789371783357: 1, 6.9720845615738405: 1, 6.970634767118219: 1, 6.966994250671041: 1, 6.965150361113812: 1, 6.945800709901644: 1, 6.937812946879745: 1, 6.937280515074351: 1, 6.933008694781485: 1, 6.929775417786: 1, 6.925648002743292: 1, 6.913834735247778: 1, 6.912625103881084: 1, 6.900227777199112: 1, 6.900009634384804: 1, 6.892943757943239: 1, 6.889800535237045: 1, 6.8836111054144435: 1, 6.878041884202959: 1, 6.867784933559784: 1, 6.863389779172354: 1, 6.854932206106438: 1, 6.852405988157565: 1, 6.846552959956564: 1, 6.844151859691189: 1, 6.841552672852482: 1, 6.838198361251475: 1, 6.837914269014631: 1, 6.827080331313124: 1, 6.826821781656983: 1, 6.821789254664449: 1, 6.816531155929945: 1, 6.812058512256804: 1, 6.806887834701513: 1, 6.799552376774305: 1, 6.7802814874270005: 1, 6.779451665205254: 1, 6.774892956484296: 1, 6.774819334543346: 1, 6.770065328553965: 1, 6.7593820153408215: 1, 6.756186914059811: 1, 6.750073670169704: 1, 6.748574675159993: 1, 6.746997794129296: 1, 6.740928443899971: 1, 6.7221963215396245: 1, 6.7212530982503536: 1, 6.706910862704476: 1, 6.70491951922496: 1, 6.693385803070201: 1, 6.683961896948499: 1, 6.680205412131914: 1, 6.672322698922836: 1, 6.665806661175904: 1, 6.661142154313514: 1, 6.65577507724585: 1, 6.649241335062643: 1, 6.648021416928578: 1, 6.641633033673139: 1, 6.628751313345468: 1, 6.628725246000161: 1, 6.624204839542654: 1, 6.621572784844526: 1, 6.605189294989049: 1, 6.597382353077788: 1, 6.597178810432583: 1, 6.593566859667499: 1, 6.592970128656291: 1, 6.586611427721229: 1, 6.572187023736954: 1, 6.568398767988924: 1, 6.565575740310147: 1, 6.560784271970288: 1, 6.5566094007898075: 1, 6.550384306477976: 1, 6.546752281871344: 1, 6.544241435805806: 1, 6.543609037480379: 1, 6.542738062694469: 1, 6.5335518452297565: 1, 6.5279738248634125: 1, 6.520731795873663: 1, 6.515663956015543: 1, 6.513465362466221: 1, 6.507397762782089: 1, 6.505858916181332: 1, 6.503620690190561: 1, 6.497602139511515: 1, 6.488910615258423: 1, 6.487489081018957: 1, 6.486721699696045: 1, 6.482267441612889: 1, 6.479197001205873: 1, 6.4738914142879835: 1, 6.472949234406288: 1, 6.4699862762332: 1, 6.467446222335095: 1, 6.463650666721645: 1, 6.459744479083084: 1, 6.452845910888976: 1, 6.446571259455178: 1, 6.437375147772991: 1, 6.434823555186791: 1, 6.428287179784081: 1, 6.427422583433763: 1, 6.426326997818246: 1, 6.425933055693564: 1, 6.423130409488598: 1, 6.416142402420761: 1, 6.41512533846111: 1, 6.407778912518549: 1, 6.4009464908265: 1, 6.400632361362674: 1, 6.399201084872064: 1, 6.398771121222739: 1, 6.398693910298181: 1, 6.395024844006849: 1, 6.392676076343998: 1, 6.38887272110073: 1, 6.379360736975719: 1, 6.379099825191223: 1, 6.377737822457731: 1, 6.372951797463017: 1, 6.369686337237917: 1, 6.362194791563312: 1, 6.36180015670087: 1, 6.361797123896236: 1, 6.3610506419622: 1, 6.359123945090875: 1, 6.35945854275128: 1, 6.349084038858023: 1, 6.348349036342279: 1, 6.344679404598878: 1, 6.344455739932798: 1, 6.34004203960962: 1, 6.338391159986534: 1, 6.337546812439763: 1, 6.336120888876584: 1, 6.331663750351538: 1, 6.329266720811574: 1, 6.328097672139696: 1, 6.321324363756329: 1, 6.31717938352454: 1, 6.30746652981277: 1, 6.302738459049485: 1, 6.292753208632827: 1, 6.28904550110206: 1, 6.288299993273792: 1, 6.283073503363987: 1, 6.28264371542526: 1, 6.28207539920521: 1, 6.28145044454061: 1, 6.280054912420066: 1, 6.277039467135268: 1, 6.275403000376772: 1, 6.269451016422114: 1, 6.267327646823247: 1, 6.2545433196258315: 1, 6.253023349923547: 1, 6.251568753226296: 1, 6.24434665718099: 1, 6.232898575683472: 1, 6.229173487082167: 1, 6.222911820131853: 1, 6.221803182144199: 1, 6.218912845628569: 1, 6.218268543010741: 1, 6.218154980443291: 1, 6.2077315422399968: 1, 6.2074901558011994: 1, 6.206919923526308: 1, 6.20364961745521: 1, 6.202828383054255: 1, 6.202559015004906: 1, 6.195229293608294: 1, 6.1872188575979905: 1, 6.183333500615305: 1, 6.178297591617514: 1, 6.176942767757961: 1, 6.174671610474641: 1, 6.169207377895481: 1, 6.168437465033292: 1, 6.157592945213799: 1, 6.157554644848806: 1, 6.155246722949498: 1, 6.152451749052985: 1, 6.1467687784834775: 1, 6.146469911774336: 1, 6.141457420800793: 1, 6.139804200996143: 1, 6.139708844559803: 1, 6.13730670595791: 1, 6.134161816682603: 1, 6.130755268268633: 1, 6.122553017970939: 1, 6.114714829553265: 1, 6.112530810578505: 1, 6.105144889714357: 1, 6.104171770556392: 1, 6.101846260640097: 1, 6.094724306815225: 1, 6.092943098216866: 1, 6.0849735700769125: 1, 6.083962496913968: 1, 6.079146713968789: 1, 6.071673934431485: 1, 6.0711453598332765: 1, 6.064758181061056: 1, 6.0634271219036044: 1, 6.062259200960211: 1, 6.061638794789515: 1, 6.058586601314483: 1, 6.0475584246199805: 1, 6.046279462568999: 1, 6.045490487100183: 1, 6.044658013657273: 1, 6.04109536460503: 1, 6.038057461514951: 1, 6.029956026292739: 1, 6.025855894704695: 1, 6.01491696631404: 1, 6.010357827805612: 1, 6.0078412390353595: 1, 6.007069750025117: 1, 6.001471769333528: 1, 6.001059607317691: 1, 5.994048838705515: 1, 5.984978566272904: 1, 5.984460451800266: 1, 5.9767807058667834: 1, 5.970776012262633: 1, 5.96795672302678: 1, 5.967845404899183: 1, 5.963400647162883: 1, 5.961562756636321: 1, 5.956777570341573: 1, 5.95338630355582: 1, 5.950583076753907: 1, 5.949754891640058: 1, 5.942235114807843: 1, 5.939590711948933: 1, 5.936765207126597: 1, 5.934913835383576: 1, 5.933172431839653: 1, 5.932564583810519: 1, 5.931169696846868: 1, 5.926717611028958: 1, 5.922843094752228: 1, 5.919945511658282: 1, 5.9156200914305765: 1, 5.905976023523879: 1, 5.904344406835265: 1, 5.903474605412722: 1, 5.902042300993354: 1, 5.9012871204115: 1, 5.899017308649157: 1, 5.893682048243449: 1, 5.893634607714694: 1, 5.89245303580313: 1, 5.887571248696374: 1, 5.885141222460431: 1, 5.8845099518488775: 1, 5.880048413319798: 1, 5.875349555408351: 1, 5.874232157114522: 1, 5.873663446397081: 1, 5.865391111762386: 1, 5.86439830344925: 1, 5.860986755122281: 1, 5.850349810698151: 1, 5.850269062533549: 1, 5.848902556404232: 1, 5.848437194185138: 1, 5.847626126763048: 1, 5.844747458973624: 1, 5.839926413051328: 1, 5.8369724795507585: 1, 5.830553868520861: 1, 5.827139530822225: 1, 5.827084330116278: 1, 5.824133798192604: 1, 5.820336453320436: 1, 5.815516377539494: 1, 5.7971343491640415: 1, 5.793743911211839: 1, 5.785615779190106: 1, 5.783934597880949: 1, 5.783656650382275: 1, 5.781304172075644: 1, 5.780564095306096: 1, 5.772070183809363: 1, 5.771436732753042: 1, 5.766227016272177: 1, 5.76237810840014: 1, 5.754913847108248: 1, 5.751524776118263: 1, 5.751132369280735: 1, 5.738782140634824: 1, 5.727508808054311: 1, 5.726199118614218: 1, 5.725805840060973: 1, 5.721699202225625: 1, 5.716566666666667: 1, 5.716566666666667: 1, 5.716566666666667: 1, 5.716566666666667: 1, 5.716566666666667: 1

5.7185066284328725: 1, 5.714866994824117: 1, 5.710035007325672: 1, 5.707796375688593: 1, 5.70682008  
5420274: 1, 5.698155250655081: 1, 5.694149505335078: 1, 5.694033274095884: 1, 5.68928421280425: 1,  
5.688401204145053: 1, 5.679812374427436: 1, 5.677170795429718: 1, 5.674814387528391: 1,  
5.670595304976204: 1, 5.661894060269134: 1, 5.661306040137116: 1, 5.65947102404416: 1,  
5.6584880932260315: 1, 5.658323917620499: 1, 5.64712457838136: 1, 5.640156218438323: 1,  
5.636375412093787: 1, 5.6322968489538: 1, 5.6261160483336035: 1, 5.620940157022657: 1,  
5.6207666909504095: 1, 5.620303810082082: 1, 5.620251550892083: 1, 5.617584858769693: 1, 5.61465357  
3975612: 1, 5.610647819508132: 1, 5.6072415958810105: 1, 5.606888893759916: 1, 5.606570509357272:  
1, 5.605240001256017: 1, 5.602445038959397: 1, 5.600876897533688: 1, 5.59110250288777: 1,  
5.589729411005058: 1, 5.5874976801355505: 1, 5.584944291052789: 1, 5.584432982554736: 1,  
5.575302203445611: 1, 5.575161135290513: 1, 5.571481858606382: 1, 5.569670139732863: 1,  
5.560665162895468: 1, 5.549313584111056: 1, 5.542468701998034: 1, 5.539425127013126: 1,  
5.53667613090923: 1, 5.530124524843596: 1, 5.523584811315713: 1, 5.520746030256738: 1,  
5.519337437814568: 1, 5.508482248199971: 1, 5.505601646972764: 1, 5.505078324501911: 1,  
5.499932497708656: 1, 5.493437571527699: 1, 5.491121846195887: 1, 5.48999850494499: 1,  
5.483296065740337: 1, 5.483267346685766: 1, 5.481951796535469: 1, 5.479222518916826: 1,  
5.478374755149457: 1, 5.477619295028376: 1, 5.47540346646769: 1, 5.471108955387635: 1,  
5.468988203354902: 1, 5.458155325970279: 1, 5.456891166237707: 1, 5.453255186404701: 1,  
5.4516544864717025: 1, 5.4510409437670715: 1, 5.448915026024611: 1, 5.4469437350279994: 1,  
5.446057819521183: 1, 5.444551474412017: 1, 5.440260453254519: 1, 5.4362938726320955: 1,  
5.4354561361046105: 1, 5.434794240876012: 1, 5.427934828979984: 1, 5.4264461336265715: 1,  
5.424238826698158: 1, 5.4236088341991975: 1, 5.422156642327516: 1, 5.410504151840336: 1,  
5.406932276044232: 1, 5.405103641465184: 1, 5.404910076233976: 1, 5.404088938669912: 1,  
5.3989351101398535: 1, 5.389629682334721: 1, 5.386141194969897: 1, 5.382292885001523: 1, 5.37974029  
8919912: 1, 5.376240024959475: 1, 5.369609144984659: 1, 5.367618660302978: 1, 5.365346559641062: 1,  
5.359978864170398: 1, 5.359815123103: 1, 5.358240313691975: 1, 5.353885098687549: 1,  
5.352413583063021: 1, 5.3464721223933145: 1, 5.344368655529633: 1, 5.341832875999072: 1,  
5.33728380588324: 1, 5.33633898722034: 1, 5.335392938151932: 1, 5.332883114125988: 1,  
5.332602432814798: 1, 5.325185380227331: 1, 5.325018419801418: 1, 5.323242590475553: 1,  
5.320100685049837: 1, 5.318292922590308: 1, 5.317758575425645: 1, 5.316370781547406: 1,  
5.310166090399736: 1, 5.30473621550924: 1, 5.298884675823079: 1, 5.297163442385255: 1,  
5.295989849359344: 1, 5.295042387572607: 1, 5.29430992952971: 1, 5.294060681689569: 1,  
5.284594989339654: 1, 5.284462124666835: 1, 5.282096467989438: 1, 5.273287407416582: 1,  
5.271778447398562: 1, 5.2609709029650515: 1, 5.258662395019501: 1, 5.25738157985273: 1,  
5.255402801130259: 1, 5.25128732227944: 1, 5.250085453300757: 1, 5.24343364969227: 1,  
5.241170839028085: 1, 5.2305189766712425: 1, 5.2297291301533475: 1, 5.221016389838603: 1,  
5.216259207081971: 1, 5.2160332773926985: 1, 5.214375713385365: 1, 5.211065083506661: 1,  
5.206089324872634: 1, 5.204776007910768: 1, 5.1986789421195585: 1, 5.192643618010112: 1,  
5.184914808589433: 1, 5.180880593026541: 1, 5.17936700557158: 1, 5.1729422470666: 1,  
5.172156757332543: 1, 5.170339584380176: 1, 5.16995742589083: 1, 5.169562420927101: 1,  
5.165090168492886: 1, 5.161395285086896: 1, 5.161029390110931: 1, 5.160780312270225: 1,  
5.149500512460107: 1, 5.147058020316481: 1, 5.140850550625411: 1, 5.1397552602721275: 1,  
5.134547622780286: 1, 5.128101806440963: 1, 5.125940578887004: 1, 5.124703943964472: 1,  
5.123316979595605: 1, 5.121158363483895: 1, 5.120012374105664: 1, 5.118766476442185: 1,  
5.117327423546183: 1, 5.113140366207468: 1, 5.112673856881091: 1, 5.110030015259387: 1,  
5.104649679598179: 1, 5.102415629460371: 1, 5.099719639496432: 1, 5.099597174688747: 1,  
5.095268846858135: 1, 5.09453714387066: 1, 5.092378342032339: 1, 5.087385041735898: 1,  
5.0799223097156805: 1, 5.078966250782268: 1, 5.0786921373866445: 1, 5.076277575632288: 1,  
5.072914563093026: 1, 5.069243001709992: 1, 5.067678458555278: 1, 5.066410964932504: 1,  
5.065035969898259: 1, 5.055050373566063: 1, 5.055035014289549: 1, 5.053365136020082: 1,  
5.050615689304438: 1, 5.049914269196854: 1, 5.045651719092932: 1, 5.04505193021577: 1,  
5.0389329801391485: 1, 5.0364147375715955: 1, 5.0355950383839945: 1, 5.027995352058877: 1,  
5.02356947487571: 1, 5.022848474753295: 1, 5.02145642526205: 1, 5.019760793532929: 1,  
5.018455751811252: 1, 5.014954366239029: 1, 5.014851543243459: 1, 5.014728395677698: 1,  
5.010372468931853: 1, 5.009666762930012: 1, 5.009373102849201: 1, 5.006196251423786: 1,  
5.00599441429096: 1, 5.0038350649570935: 1, 5.0023377491675936: 1, 5.000215836240911: 1,  
4.996290935749122: 1, 4.99535879735655: 1, 4.995116222898656: 1, 4.994755259290891: 1,  
4.994415327706056: 1, 4.993650768433: 1, 4.991294212406486: 1, 4.9868875906302295: 1,  
4.986138447821082: 1, 4.984757214109896: 1, 4.983108105475257: 1, 4.98000296289694: 1,  
4.9795661983313035: 1, 4.977953079073226: 1, 4.977417979939025: 1, 4.976397574524791: 1, 4.97282642  
8212472: 1, 4.970241877603853: 1, 4.968058371799982: 1, 4.963146894202303: 1, 4.963018656301361: 1,  
4.962686593455276: 1, 4.962200105432318: 1, 4.946880279918504: 1, 4.945588978160046: 1,  
4.941625103660423: 1, 4.940453194126887: 1, 4.936308822455724: 1, 4.9232957454430935: 1,  
4.915801189393756: 1, 4.913116685608693: 1, 4.909966158516649: 1, 4.908845435872575: 1,  
4.908366851388559: 1, 4.897748233170453: 1, 4.896240959847156: 1, 4.895909578497357: 1,  
4.895118675040717: 1, 4.895062360427643: 1, 4.892730627063446: 1, 4.890119678025315: 1,  
4.887848089803108: 1, 4.887336907339567: 1, 4.885434911105181: 1, 4.878587497722377: 1,  
4.874763535346076: 1, 4.871171672868576: 1, 4.8656238105499465: 1, 4.864517815143099: 1,  
4.8643581692778515: 1, 4.859203056522621: 1, 4.858521000159169: 1, 4.857882251859499: 1, 4.85229258  
88837545: 1, 4.8492460972364: 1, 4.849040112627673: 1, 4.840294832846789: 1, 4.839705898747609: 1,  
4.834377168837295: 1, 4.831444734298727: 1, 4.831196677636456: 1, 4.830857305624724: 1,  
4.830612189644915: 1, 4.83035443485908: 1, 4.828364848354913: 1, 4.8268869300088575: 1,  
4.821530510588967: 1, 4.817824980421616: 1, 4.816884429485378: 1, 4.8157413810422005: 1,  
4.810807211798044: 1, 4.81069205019542: 1, 4.80990684521169: 1, 4.808992246644466: 1,  
4.805910104604024: 1, 4.802363407856109: 1, 4.798104981492134: 1, 4.795855700662849: 1,  
4.793053461843088: 1, 4.792867336647074: 1, 4.789337463139297: 1, 4.7883044136941555: 1,  
4.776551219688534: 1, 4.775750469713873: 1, 4.774205325233771: 1, 4.773963401673264: 1,  
4.773750469713873: 1, 4.773750469713873: 1, 4.773750469713873: 1, 4.773750469713873: 1

4.773307015431376: 1, 4.772178503115321: 1, 4.770115806228567: 1, 4.769631640397137: 1, 4.765470730073898: 1, 4.765323043978703: 1, 4.762937310180706: 1, 4.758819112660872: 1, 4.75187341881589: 1, 4.748238522748845: 1, 4.745161121192289: 1, 4.742054382834529: 1, 4.741615241733631: 1, 4.740933806716363: 1, 4.739490755289273: 1, 4.738936051242229: 1, 4.737844173082161: 1, 4.7352988867829024: 1, 4.73523693138319: 1, 4.734370645465436: 1, 4.734020048606232: 1, 4.72047401996755: 1, 4.71825250437165: 1, 4.717314436248292: 1, 4.7150519384145575: 1, 4.711757150397257: 1, 4.710864874887621: 1, 4.707829123807485: 1, 4.70378823 32746355: 1, 4.70296921643569: 1, 4.702157307351482: 1, 4.693075005854421: 1, 4.6872592831614694: 1, 4.686823849845743: 1, 4.684584049283698: 1, 4.6818815846128885: 1, 4.679587520287394: 1, 4.677595101978762: 1, 4.676674489725872: 1, 4.670291203486952: 1, 4.666188617472765: 1, 4.664623911345924: 1, 4.657243926055926: 1, 4.656982129026654: 1, 4.651561795506981: 1, 4.650378036094082: 1, 4.6493389441419115: 1, 4.647823175941512: 1, 4.646187060000831: 1, 4.643619505579293: 1, 4.640641258225292: 1, 4.6395596539649455: 1, 4.63915612269569: 1, 4.638190700444934: 1, 4.636331710761858: 1, 4.634172507542458: 1, 4.634010174499257: 1, 4.630426555114664: 1, 4.626713227302652: 1, 4.625716644775695: 1, 4.625555011742524: 1, 4.621323609255618: 1, 4.61743323840002: 1, 4.617042018310118: 1, 4.614775534363596: 1, 4.61466873529294: 1, 4.612394699928937: 1, 4.612180167649222: 1, 4.610770302233104: 1, 4.6062550192782465: 1, 4.602614934683913: 1, 4.596551850401169: 1, 4.585277336943502: 1, 4.57277146 7194507: 1, 4.568096037965321: 1, 4.566969588336921: 1, 4.5651256293232585: 1, 4.553635101081163: 1, 4.546726442282488: 1, 4.5427948547986805: 1, 4.542020036091564: 1, 4.532957621616767: 1, 4.531606025069032: 1, 4.53153322842507: 1, 4.527907557253492: 1, 4.516250180769683: 1, 4.515295884473163: 1, 4.5141435068468425: 1, 4.512155898977833: 1, 4.500431686355015: 1, 4.492341218304294: 1, 4.476359803510155: 1, 4.4716711591295875: 1, 4.470177469685173: 1, 4.4691235540951775: 1, 4.464560917747493: 1, 4.461967256360174: 1, 4.461731036909922: 1, 4.45735284 3525774: 1, 4.453843014449613: 1, 4.451030837865101: 1, 4.447853801776508: 1, 4.444206905483644: 1, 4.438206658506043: 1, 4.434971828729152: 1, 4.431738205385265: 1, 4.426284648508997: 1, 4.419111294964117: 1, 4.418597905365838: 1, 4.410981971150588: 1, 4.408569180618263: 1, 4.40442235641261: 1, 4.401970457316341: 1, 4.401604808289976: 1, 4.401565442514928: 1, 4.395361338169909: 1, 4.394713528785653: 1, 4.394275830931278: 1, 4.389213740717272: 1, 4.384290647887208: 1, 4.382733291590986: 1, 4.38062237969791: 1, 4.379165809537193: 1, 4.371037950098307: 1, 4.368725023598987: 1, 4.368702490369654: 1, 4.357127030276572: 1, 4.357035448076488: 1, 4.351495929590222: 1, 4.342898650991759: 1, 4.3408380779212745: 1, 4.340586723898401: 1, 4.340181816582982: 1, 4.339139654629641: 1, 4.333778981699127: 1, 4.332863078417063: 1, 4.331727871925859: 1, 4.329496243404825: 1, 4.329152811228107: 1, 4.3270572740381015: 1, 4.326513679594772: 1, 4.320791810979486: 1, 4.320243092081452: 1, 4.31394910 9154409: 1, 4.313522892764479: 1, 4.300910717847041: 1, 4.2914070633842: 1, 4.289236751603193: 1, 4.288760141360002: 1, 4.287463286382373: 1, 4.286783586922946: 1, 4.283656114594232: 1, 4.276240256095293: 1, 4.269436982982311: 1, 4.26475261397136: 1, 4.264549584636356: 1, 4.262872727482005: 1, 4.251842275151268: 1, 4.251789003017001: 1, 4.250417247678242: 1, 4.246639596772416: 1, 4.241457212168962: 1, 4.237539151848959: 1, 4.231866965161524: 1, 4.2275350931562645: 1, 4.220146768915774: 1, 4.217746106416265: 1, 4.210870600105587: 1, 4.20995296 4654822: 1, 4.208879730239735: 1, 4.20854453960697: 1, 4.207450649268616: 1, 4.2050634888748535: 1, 4.200990922493716: 1, 4.1970771151478745: 1, 4.194773342302923: 1, 4.193919355612454: 1, 4.192744670594784: 1, 4.18960530244654: 1, 4.1855036023979055: 1, 4.182588832607453: 1, 4.18043934906325: 1, 4.178378575183917: 1, 4.175973690530259: 1, 4.171020184925834: 1, 4.16520603271201: 1, 4.163928474629289: 1, 4.160520032244243: 1, 4.1546208755410022: 1, 4.139176248665656: 1, 4.136636035860688: 1, 4.1330530786337265: 1, 4.132343823784156: 1, 4.129203996357184: 1, 4.1283753198165165: 1, 4.120982063799175: 1, 4.117788984995421: 1, 4.104316087975984: 1, 4.103041260074071: 1, 4.101195755675666: 1, 4.099539019751617: 1, 4.098479745922179: 1, 4.093376723119562: 1, 4.085044747460734: 1, 4.0744788769750935: 1, 4.065972686821436: 1, 4.0657041075642235: 1, 4.062644737873088: 1, 4.060193023315892: 1, 4.050886351265687: 1, 4.050159722476665: 1, 4.039995087653942: 1, 4.035713330863368: 1, 4.033263730446443: 1, 4.03324455300392: 1, 4.021842414484016: 1, 4.021654206908767: 1, 4.020864442184634: 1, 4.014647112243779: 1, 4.010288400742972: 1, 4.004939326369199: 1, 3.999234462076511: 1, 3.9787652351069664: 1, 3.978282093463443: 1, 3.9641905041084953: 1, 3.9615310912263157: 1, 3.954999366794561: 1, 3.9494876396556804: 1, 3.930910543091535: 1, 3.9265581952217588: 1, 3.9170056086488567: 1, 3.9159189329310498: 1, 3.913811001911172: 1, 3.9016536056097277: 1, 3.901602663344786: 1, 3.892639367424239: 1, 3.888479163995114: 1, 3.88677746 98041235: 1, 3.8859530995501492: 1, 3.8833757298830807: 1, 3.8643280036820573: 1, 3.862595799614239: 1, 3.856991673338055: 1, 3.8553534291677627: 1, 3.8304218564217702: 1, 3.8140695508338895: 1, 3.784368559937053: 1, 3.7818876841223217: 1, 3.766212863554476: 1, 3.76008440675485: 1, 3.7579953737410645: 1, 3.7554984336773614: 1, 3.737079753367513: 1, 3.731023062501971: 1, 3.714630209528808: 1, 3.7113095345991995: 1, 3.703335940592092: 1, 3.6991104191987314: 1, 3.692806724228393: 1, 3.685041005469602: 1, 3.6800714948941584: 1, 3.6670021103478807: 1, 3.638773654221167: 1, 3.6307218490605943: 1, 3.5995065962705493: 1, 3.58925970646872: 1, 3.561762066825263: 1, 3.5527486298816426: 1, 3.5315851225603274: 1, 3.5300195765981393: 1, 3.4315456701777634: 1, 3.348318957722383: 1, 3.1693411128895326: 1}})

In [47]:

```
# Train a Logistic regression+Calibration model using text features which are one-hot encoded
alpha = [10 ** x for x in range(-5, 1)]

# read more about SGDClassifier() at http://scikit-learn.org/stable/modules/generated/sklearn.linear_model.SGDClassifier.html
# -----
```

```

# default parameters
# SGDClassifier(loss='hinge', penalty='l2', alpha=0.0001, l1_ratio=0.15, fit_intercept=True, max_iter=None, tol=None,
# shuffle=True, verbose=0, epsilon=0.1, n_jobs=1, random_state=None, learning_rate='optimal', eta0
# =0.0, power_t=0.5,
# class_weight=None, warm_start=False, average=False, n_iter=None)

# some of methods
# fit(X, y[, coef_init, intercept_init, ...]) Fit linear model with Stochastic Gradient Descent.
# predict(X) Predict class labels for samples in X.

#-----
# video link:
#-----

cv_log_error_array=[]
for i in alpha:
    clf = SGDClassifier(alpha=i, penalty='l2', loss='log', random_state=42)
    clf.fit(train_text_feature_onehotCoding, y_train)

    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_text_feature_onehotCoding, y_train)
    predict_y = sig_clf.predict_proba(cv_text_feature_onehotCoding)
    cv_log_error_array.append(log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))
    print('For values of alpha = ', i, "The log loss is:", log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))


fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array, c='g')
for i, txt in enumerate(np.round(cv_log_error_array, 3)):
    ax.annotate((alpha[i], np.round(txt, 3)), (alpha[i], cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

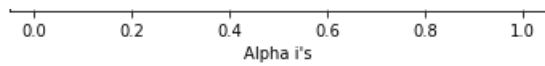
best_alpha = np.argmin(cv_log_error_array)
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=42)
clf.fit(train_text_feature_onehotCoding, y_train)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_text_feature_onehotCoding, y_train)

predict_y = sig_clf.predict_proba(train_text_feature_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:", log_loss(y_train, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(cv_text_feature_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The cross validation log loss is:", log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(test_text_feature_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:", log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))

```

For values of alpha = 1e-05 The log loss is: 1.140779489415932  
 For values of alpha = 0.0001 The log loss is: 1.0758027258362972  
 For values of alpha = 0.001 The log loss is: 1.29578788031302  
 For values of alpha = 0.01 The log loss is: 1.6207186003778338  
 For values of alpha = 0.1 The log loss is: 1.9185358313165983  
 For values of alpha = 1 The log loss is: 1.934455072880182





For values of best alpha = 0.0001 The train log loss is: 0.7360639947052967  
For values of best alpha = 0.0001 The cross validation log loss is: 1.0758027258362972  
For values of best alpha = 0.0001 The test log loss is: 1.2094711851133668

**Q. Is the Text feature stable across all the data sets (Test, Train, Cross validation)?**

**Ans. Yes, it seems like!**

In [48]:

```
def get_intersec_text(df):
    df_text_vec = CountVectorizer(min_df=3)
    df_text_fea = df_text_vec.fit_transform(df['TEXT'])
    df_text_features = df_text_vec.get_feature_names()

    df_text_fea_counts = df_text_fea.sum(axis=0).A1
    df_text_fea_dict = dict(zip(list(df_text_features), df_text_fea_counts))
    len1 = len(set(df_text_features))
    len2 = len(set(train_text_features) & set(df_text_features))
    return len1, len2
```

In [49]:

```
len1, len2 = get_intersec_text(test_df)
print(np.round((len2/len1)*100, 3), "% of word of test data appeared in train data")
len1, len2 = get_intersec_text(cv_df)
print(np.round((len2/len1)*100, 3), "% of word of Cross Validation appeared in train data")
```

5.806 % of word of test data appeared in train data  
6.618 % of word of Cross Validation appeared in train data

## 4. Machine Learning Models

In [50]:

```
#Data preparation for ML models.

#Misc. functionns for ML models

def predict_and_plot_confusion_matrix(train_x, train_y, test_x, test_y, clf):
    clf.fit(train_x, train_y)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_x, train_y)
    pred_y = sig_clf.predict(test_x)

    # for calculating log_loss we will provide the array of probabilities belongs to each class
    print("Log loss :", log_loss(test_y, sig_clf.predict_proba(test_x)))
    # calculating the number of data points that are misclassified
    print("Number of mis-classified points :", np.count_nonzero((pred_y - test_y))/test_y.shape[0])
    plot_confusion_matrix(test_y, pred_y)
```

In [51]:

```
def report_log_loss(train_x, train_y, test_x, test_y, clf):
    clf.fit(train_x, train_y)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_x, train_y)
    sig_clf_probs = sig_clf.predict_proba(test_x)
    return log_loss(test_y, sig_clf_probs, eps=1e-15)
```

In [52]:

```
# this function will be used just for naive bayes
# for the given indices, we will print the name of the features
```



```

# and we will check whether the feature present in the test point text or not
def get_impfeature_names(indices, text, gene, var, no_features):
    gene_count_vec = CountVectorizer()
    var_count_vec = CountVectorizer()
    text_count_vec = CountVectorizer(min_df=3)

    gene_vec = gene_count_vec.fit(train_df['Gene'])
    var_vec = var_count_vec.fit(train_df['Variation'])
    text_vec = text_count_vec.fit(train_df['TEXT'])

    fea1_len = len(gene_vec.get_feature_names())
    fea2_len = len(var_count_vec.get_feature_names())

    word_present = 0
    for i,v in enumerate(indices):
        if (v < fea1_len):
            word = gene_vec.get_feature_names()[v]
            yes_no = True if word == gene else False
            if yes_no:
                word_present += 1
                print(i, "Gene feature [{}] present in test data point [{}].format(word,yes_no))
        elif (v < fea1_len+fea2_len):
            word = var_vec.get_feature_names()[v-(fea1_len)]
            yes_no = True if word == var else False
            if yes_no:
                word_present += 1
                print(i, "variation feature [{}] present in test data point [{}].format(word,yes_r
o))
        else:
            word = text_vec.get_feature_names()[v-(fea1_len+fea2_len)]
            yes_no = True if word in text.split() else False
            if yes_no:
                word_present += 1
                print(i, "Text feature [{}] present in test data point [{}].format(word,yes_no))

    print("Out of the top ",no_features," features ", word_present, "are present in query point")

```

## Stacking the three types of features

In [53]:

```

# merging gene, variance and text features

# building train, test and cross validation data sets
# a = [[1, 2],
#       [3, 4]]
# b = [[4, 5],
#       [6, 7]]
# hstack(a, b) = [[1, 2, 4, 5],
#                 [ 3, 4, 6, 7]]

train_gene_var_onehotCoding =
hstack((train_gene_feature_onehotCoding,train_variation_feature_onehotCoding))
test_gene_var_onehotCoding =
hstack((test_gene_feature_onehotCoding,test_variation_feature_onehotCoding))
cv_gene_var_onehotCoding = hstack((cv_gene_feature_onehotCoding,cv_variation_feature_onehotCoding)
)

train_x_onehotCoding = hstack((train_gene_var_onehotCoding, train_text_feature_onehotCoding)).tocsr()
train_y = np.array(list(train_df['Class']))

test_x_onehotCoding = hstack((test_gene_var_onehotCoding, test_text_feature_onehotCoding)).tocsr()
test_y = np.array(list(test_df['Class']))

cv_x_onehotCoding = hstack((cv_gene_var_onehotCoding, cv_text_feature_onehotCoding)).tocsr()
cv_y = np.array(list(cv_df['Class']))

train_gene_var_responseCoding =
np.hstack((train_gene_feature_responseCoding,train_variation_feature_responseCoding))
test_gene_var_responseCoding =
np.hstack((test_gene_feature_responseCoding,test_variation_feature_responseCoding))

```

```

cv_gene_var_responseCoding =
np.hstack((cv_gene_feature_responseCoding,cv_variation_feature_responseCoding))

train_x_responseCoding = np.hstack((train_gene_var_responseCoding,
train_text_feature_responseCoding))
test_x_responseCoding = np.hstack((test_gene_var_responseCoding, test_text_feature_responseCoding)
)
cv_x_responseCoding = np.hstack((cv_gene_var_responseCoding, cv_text_feature_responseCoding))

```

In [54]:

```

print("One hot encoding features :")
print("(number of data points * number of features) in train data = ", train_x_onehotCoding.shape)
print("(number of data points * number of features) in test data = ", test_x_onehotCoding.shape)
print("(number of data points * number of features) in cross validation data =", cv_x_onehotCoding
.shape)

```

One hot encoding features :

(number of data points \* number of features) in train data = (2124, 4226)

(number of data points \* number of features) in test data = (665, 4226)

(number of data points \* number of features) in cross validation data = (532, 4226)

In [55]:

```

print(" Response encoding features :")
print("(number of data points * number of features) in train data = ", train_x_responseCoding.shap
e)
print("(number of data points * number of features) in test data = ", test_x_responseCoding.shape)
print("(number of data points * number of features) in cross validation data =",
cv_x_responseCoding.shape)

```

Response encoding features :

(number of data points \* number of features) in train data = (2124, 27)

(number of data points \* number of features) in test data = (665, 27)

(number of data points \* number of features) in cross validation data = (532, 27)

## 4.1. Base Line Model

### 4.1.1. Naive Bayes

#### 4.1.1.1. Hyper parameter tuning

In [56]:

```

# find more about Multinomial Naive base function here http://scikit-learn.org/stable/modules/generated/sklearn.naive\_bayes.MultinomialNB.html
# -----
# default paramters
# sklearn.naive_bayes.MultinomialNB(alpha=1.0, fit_prior=True, class_prior=None)

# some of methods of MultinomialNB()
# fit(X, y[, sample_weight]) Fit Naive Bayes classifier according to X, y
# predict(X) Perform classification on an array of test vectors X.
# predict_log_proba(X) Return log-probability estimates for the test vector X.
# -----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/naive-bayes-algorithm-1/
# -----

# find more about CalibratedClassifierCV here at http://scikit-learn.org/stable/modules/generated/sklearn.calibration.CalibratedClassifierCV.html
# -----
# default paramters
# sklearn.calibration.CalibratedClassifierCV(base_estimator=None, method='sigmoid', cv=3)
#
# some of the methods of CalibratedClassifierCV()
# fit(X, y[, sample_weight]) Fit the calibrated model
# get_params([deep]) Get parameters for this estimator.

```

```

# predict(X) Predict the target of new samples.
# predict_proba(X) Posterior probabilities of classification
# -----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/naive-bayes-algorithm-1/
# -----

alpha = [0.00001, 0.0001, 0.001, 0.1, 1, 10, 100, 1000]
cv_log_error_array = []
for i in alpha:
    print("for alpha =", i)
    clf = MultinomialNB(alpha=i)
    clf.fit(train_x_onehotCoding, train_y)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_x_onehotCoding, train_y)
    sig_clf_probs = sig_clf.predict_proba(cv_x_onehotCoding)
    cv_log_error_array.append(log_loss(cv_y, sig_clf_probs, labels=clf.classes_, eps=1e-15))
    # to avoid rounding error while multiplying probabilities we use log-probability estimates
    print("Log Loss :", log_loss(cv_y, sig_clf_probs))

fig, ax = plt.subplots()
ax.plot(np.log10(alpha), cv_log_error_array, c='g')
for i, txt in enumerate(np.round(cv_log_error_array, 3)):
    ax.annotate((alpha[i], str(txt)), (np.log10(alpha[i]), cv_log_error_array[i]))
plt.grid()
plt.xticks(np.log10(alpha))
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

best_alpha = np.argmin(cv_log_error_array)
clf = MultinomialNB(alpha=alpha[best_alpha])
clf.fit(train_x_onehotCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_onehotCoding, train_y)


predict_y = sig_clf.predict_proba(train_x_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:", log_loss(y_train,
predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(cv_x_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The cross validation log loss is:", log_lo
ss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(test_x_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:", log_loss(y_test, p
redict_y, labels=clf.classes_, eps=1e-15))


```

```

for alpha = 1e-05
Log Loss : 1.153105050504226
for alpha = 0.0001
Log Loss : 1.152707984150967
for alpha = 0.001
Log Loss : 1.1504317324504902
for alpha = 0.1
Log Loss : 1.199030089131308
for alpha = 1
Log Loss : 1.2969612308356298
for alpha = 10
Log Loss : 1.4827173823549864
for alpha = 100
Log Loss : 1.4525345412018227
for alpha = 1000
Log Loss : 1.4266255568324775

```





For values of best alpha = 0.001 The train log loss is: 0.5812146150905396  
 For values of best alpha = 0.001 The cross validation log loss is: 1.1504317324504902  
 For values of best alpha = 0.001 The test log loss is: 1.2333635964950285

#### 4.1.1.2. Testing the model with best hyper paramters

In [57]:

```
# find more about Multinomial Naive base function here http://scikit-learn.org/stable/modules/generated/sklearn.naive\_bayes.MultinomialNB.html
# -----
# default paramters
# sklearn.naive_bayes.MultinomialNB(alpha=1.0, fit_prior=True, class_prior=None)

# some of methods of MultinomialNB()
# fit(X, y[, sample_weight]) Fit Naive Bayes classifier according to X, y
# predict(X) Perform classification on an array of test vectors X.
# predict_log_proba(X) Return log-probability estimates for the test vector X.
# -----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/naive-bayes-algorithm-1/
# -----


# find more about CalibratedClassifierCV here at http://scikit-learn.org/stable/modules/generated/sklearn.calibration.CalibratedClassifierCV.html
# -----
# default paramters
# sklearn.calibration.CalibratedClassifierCV(base_estimator=None, method='sigmoid', cv=3)
#
# some of the methods of CalibratedClassifierCV()
# fit(X, y[, sample_weight]) Fit the calibrated model
# get_params([deep]) Get parameters for this estimator.
# predict(X) Predict the target of new samples.
# predict_proba(X) Posterior probabilities of classification
# -----


clf = MultinomialNB(alpha=alpha[best_alpha])
clf.fit(train_x_onehotCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_onehotCoding, train_y)
sig_clf_probs = sig_clf.predict_proba(cv_x_onehotCoding)
# to avoid rounding error while multiplying probabilites we use log-probability estimates
print("Log Loss :", log_loss(cv_y, sig_clf_probs))
print("Number of missclassified point :", np.count_nonzero((sig_clf.predict(cv_x_onehotCoding) - cv_y) / cv_y.shape[0]))
plot_confusion_matrix(cv_y, sig_clf.predict(cv_x_onehotCoding.toarray()))
```

Log Loss : 1.1504317324504902


Number of missclassified point : 0.37593984962406013

----- Confusion matrix -----






----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----



#### 4.1.1.3. Feature Importance, Correctly classified point

In [58]:

```
test_point_index = 1
no_feature = 100
predicted_cls = sig_clf.predict(test_x_onehotCoding[test_point_index])
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:",
      np.round(sig_clf.predict_proba(test_x_onehotCoding[test_point_index]),4))
print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-clf.coef_)[predicted_cls-1][:,no_feature]
print("-"*50)
get_impfeature_names(indices[0],
test_df['TEXT'].iloc[test_point_index],test_df['Gene'].iloc[test_point_index],test_df['Variation']
.iloc[test_point_index], no_feature)
```

```

Predicted Class : 1
Predicted Class Probabilities: [[0.462  0.0468 0.0116 0.3135 0.0379 0.0376 0.0842 0.0041 0.0021]]
Actual Class : 3
-----
48 Text feature [15] present in test data point [True]
88 Text feature [185] present in test data point [True]
Out of the top 100 features 2 are present in query point

```

#### 4.1.1.4. Feature Importance, Incorrectly classified point

In [59]:

```

test_point_index = 100
no_feature = 100
predicted_cls = sig_clf.predict(test_x_onehotCoding[test_point_index])
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:",
      np.round(sig_clf.predict_proba(test_x_onehotCoding[test_point_index]),4))
print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-clf.coef_) [predicted_cls-1][:,no_feature]
print("-"*50)
get_impfeature_names(indices[0],
test_df['TEXT'].iloc[test_point_index],test_df['Gene'].iloc[test_point_index],test_df['Variation']
.iloc[test_point_index], no_feature)

```

```

Predicted Class : 1
Predicted Class Probabilities: [[0.5557 0.0517 0.0128 0.197  0.0416 0.0416 0.0928 0.0045 0.0023]]
Actual Class : 4
-----
48 Text feature [15] present in test data point [True]
Out of the top 100 features 1 are present in query point

```

## 4.2. K Nearest Neighbour Classification

### 4.2.1. Hyper parameter tuning

In [60]:

```

# find more about KNeighborsClassifier() here http://scikit-
learn.org/stable/modules/generated/sklearn.neighbors.KNeighborsClassifier.html
# -----
# default parameter
# KNeighborsClassifier(n_neighbors=5, weights='uniform', algorithm='auto', leaf_size=30, p=2,
# metric='minkowski', metric_params=None, n_jobs=1, **kwargs)

# methods of
# fit(X, y) : Fit the model using X as training data and y as target values
# predict(X):Predict the class labels for the provided data
# predict_proba(X):Return probability estimates for the test data X.
#-----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/k-nearest-ne
ighbors-geometric-intuition-with-a-toy-example-1/
#-----

# find more about CalibratedClassifierCV here at http://scikit-
learn.org/stable/modules/generated/sklearn.calibration.CalibratedClassifierCV.html
# -----
# default paramters
# sklearn.calibration.CalibratedClassifierCV(base_estimator=None, method='sigmoid', cv=3)
#
# some of the methods of CalibratedClassifierCV()
# fit(X, y[, sample_weight]) Fit the calibrated model
# get_params([deep]) Get parameters for this estimator.
# predict(X) Predict the target of new samples.
# predict_proba(X) Posterior probabilities of classification
#-----
# video link:
#-----

```

```

alpha = [5, 11, 15, 21, 31, 41, 51, 99]
cv_log_error_array = []
for i in alpha:
    print("for alpha =", i)
    clf = KNeighborsClassifier(n_neighbors=i)
    clf.fit(train_x_responseCoding, train_y)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_x_responseCoding, train_y)
    sig_clf_probs = sig_clf.predict_proba(cv_x_responseCoding)
    cv_log_error_array.append(log_loss(cv_y, sig_clf_probs, labels=clf.classes_, eps=1e-15))
    # to avoid rounding error while multiplying probabilities we use log-probability estimates
    print("Log Loss :", log_loss(cv_y, sig_clf_probs))

fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array, c='g')
for i, txt in enumerate(np.round(cv_log_error_array, 3)):
    ax.annotate((alpha[i], str(txt)), (alpha[i], cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

best_alpha = np.argmin(cv_log_error_array)
clf = KNeighborsClassifier(n_neighbors=alpha[best_alpha])
clf.fit(train_x_responseCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_responseCoding, train_y)


predict_y = sig_clf.predict_proba(train_x_responseCoding)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:", log_loss(y_train,
predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(cv_x_responseCoding)
print('For values of best alpha = ', alpha[best_alpha], "The cross validation log loss is:", log_lo
ss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(test_x_responseCoding)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:", log_loss(y_test, p
redict_y, labels=clf.classes_, eps=1e-15))

```

```

for alpha = 5
Log Loss : 1.0147585241733181
for alpha = 11
Log Loss : 1.0100418537371045
for alpha = 15
Log Loss : 1.017000427290506
for alpha = 21
Log Loss : 1.020051465209906
for alpha = 31
Log Loss : 1.0425869526802527
for alpha = 41
Log Loss : 1.0358922726241941
for alpha = 51
Log Loss : 1.0312499090574307
for alpha = 99
Log Loss : 1.0316815087635784

```



For values of best alpha = 11 The train log loss is: 0.6614049001535538

For values of best alpha = 11 The cross validation log loss is: 1.0100418537371045  
 For values of best alpha = 11 The test log loss is: 1.0813518997758362

## 4.2.2. Testing the model with best hyper paramters

In [61]:


```
# find more about KNeighborsClassifier() here http://scikit-learn.org/stable/modules/generated/sklearn.neighbors.KNeighborsClassifier.html
# -----
# default parameter
# KNeighborsClassifier(n_neighbors=5, weights='uniform', algorithm='auto', leaf_size=30, p=2,
# metric='minkowski', metric_params=None, n_jobs=1, **kwargs)

# methods of
# fit(X, y) : Fit the model using X as training data and y as target values
# predict(X):Predict the class labels for the provided data
# predict_proba(X):Return probability estimates for the test data X.
#-----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/k-nearest-neighbors-geometric-intuition-with-a-toy-example-1/
#-----
clf = KNeighborsClassifier(n_neighbors=alpha[best_alpha])
predict_and_plot_confusion_matrix(train_x_responseCoding, train_y, cv_x_responseCoding, cv_y, clf)
```


Log loss : 1.0100418537371045

Number of mis-classified points : 0.3533834586466165

----- Confusion matrix -----




----- Precision matrix (Column Sum=1) -----





----- Recall matrix (Row sum=1) -----



### 4.2.3. Sample Query point -1

In [62]:

```
clf = KNeighborsClassifier(n_neighbors=alpha[best_alpha])
clf.fit(train_x_responseCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_responseCoding, train_y)

test_point_index = 1
predicted_cls = sig_clf.predict(test_x_responseCoding[0].reshape(1,-1))
print("Predicted Class :", predicted_cls[0])
print("Actual Class :", test_y[test_point_index])
neighbors = clf.kneighbors(test_x_responseCoding[test_point_index].reshape(1, -1), alpha[best_alpha])
print("The ", alpha[best_alpha], " nearest neighbours of the test points belongs to classes", train_y[neighbors[1][0]])
print("Frequency of nearest points :", Counter(train_y[neighbors[1][0]]))
```

Predicted Class : 7

Actual Class : 3

The 11 nearest neighbours of the test points belongs to classes [1 4 4 1 4 1 1 4 4 4 1]

Frequency of nearest points : Counter({4: 6, 1: 5})

### 4.2.4. Sample Query Point-2

In [63]:

```
clf = KNeighborsClassifier(n_neighbors=alpha[best_alpha])
clf.fit(train_x_responseCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_responseCoding, train_y)

test_point_index = 100

predicted_cls = sig_clf.predict(test_x_responseCoding[test_point_index].reshape(1,-1))
print("Predicted Class :", predicted_cls[0])
print("Actual Class :", test_y[test_point_index])
neighbors = clf.kneighbors(test_x_responseCoding[test_point_index].reshape(1, -1), alpha[best_alpha])
print("the k value for knn is", alpha[best_alpha], "and the nearest neighbours of the test points belongs to classes", train_y[neighbors[1][0]])
print("Frequency of nearest points :", Counter(train_y[neighbors[1][0]]))
```

Predicted Class : 4

Actual Class : 4

the k value for knn is 11 and the nearest neighbours of the test points belongs to classes [4 1 1  
5 1 4 4 4 4 4 1]  
Frequency of nearest points : Counter({4: 6, 1: 4, 5: 1})

## 4.3. Logistic Regression

### 4.3.1. With Class balancing

#### 4.3.1.1. Hyper paramter tuning

In [64]:

```
# read more about SGDClassifier() at http://scikit-learn.org/stable/modules/generated/sklearn.linear\_model.SGDClassifier.html
# -----
# default parameters
# SGDClassifier(loss='hinge', penalty='l2', alpha=0.0001, l1_ratio=0.15, fit_intercept=True, max_iter=None, tol=None,
# shuffle=True, verbose=0, epsilon=0.1, n_jobs=1, random_state=None, learning_rate='optimal', eta0=0.0, power_t=0.5,
# class_weight=None, warm_start=False, average=False, n_iter=None)

# some of methods
# fit(X, y[, coef_init, intercept_init, ...]) Fit linear model with Stochastic Gradient Descent.
# predict(X) Predict class labels for samples in X.

#-----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/geometric-in-tuition-1/
#-----

# find more about CalibratedClassifierCV here at http://scikit-learn.org/stable/modules/generated/sklearn.calibration.CalibratedClassifierCV.html
# -----
# default paramters
# sklearn.calibration.CalibratedClassifierCV(base_estimator=None, method='sigmoid', cv=3)
#
# some of the methods of CalibratedClassifierCV()
# fit(X, y[, sample_weight]) Fit the calibrated model
# get_params([deep]) Get parameters for this estimator.
# predict(X) Predict the target of new samples.
# predict_proba(X) Posterior probabilities of classification
#-----
# video link:
#-----

alpha = [10 ** x for x in range(-6, 3)]
cv_log_error_array = []
for i in alpha:
    print("for alpha =", i)
    clf = SGDClassifier(class_weight='balanced', alpha=i, penalty='l2', loss='log', random_state=42)

    clf.fit(train_x_onehotCoding, train_y)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_x_onehotCoding, train_y)
    sig_clf_probs = sig_clf.predict_proba(cv_x_onehotCoding)
    cv_log_error_array.append(log_loss(cv_y, sig_clf_probs, labels=clf.classes_, eps=1e-15))
    # to avoid rounding error while multiplying probabilitites we use log-probability estimates
    print("Log Loss :", log_loss(cv_y, sig_clf_probs))

fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array, c='g')
for i, txt in enumerate(np.round(cv_log_error_array, 3)):
    ax.annotate((alpha[i], str(txt)), (alpha[i], cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()
```

```

best_alpha = np.argmin(cv_log_error_array)
clf = SGDClassifier(class_weight='balanced', alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=42)
clf.fit(train_x_onehotCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_onehotCoding, train_y)


predict_y = sig_clf.predict_proba(train_x_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:", log_loss(y_train, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(cv_x_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The cross validation log loss is:", log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(test_x_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:", log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))

```

```

for alpha = 1e-06
Log Loss : 1.0270228120402658
for alpha = 1e-05
Log Loss : 1.0172016464993268
for alpha = 0.0001
Log Loss : 0.9592513281465967
for alpha = 0.001
Log Loss : 0.9554926080094617
for alpha = 0.01
Log Loss : 1.1118829662812784
for alpha = 0.1
Log Loss : 1.6447457229985016
for alpha = 1
Log Loss : 1.793442602408434
for alpha = 10
Log Loss : 1.8090578411240252
for alpha = 100
Log Loss : 1.8108508261835534

```



```

For values of best alpha = 0.001 The train log loss is: 0.6779496215824264
For values of best alpha = 0.001 The cross validation log loss is: 0.9554926080094617
For values of best alpha = 0.001 The test log loss is: 1.087019650771065

```

#### 4.3.1.2. Testing the model with best hyper parameters

In [65]:

```

# read more about SGDClassifier() at http://scikit-learn.org/stable/modules/generated/sklearn.linear_model.SGDClassifier.html
# -----
# default parameters
# SGDClassifier(loss='hinge', penalty='l2', alpha=0.0001, l1_ratio=0.15, fit_intercept=True, max_iter=None, tol=None,
# shuffle=True, verbose=0, epsilon=0.1, n_jobs=1, random_state=None, learning_rate='optimal', eta0=0.0, power_t=0.5,
# class_weight=None, warm_start=False, average=False, n_iter=None)

# some of methods
# fit(X, y[, coef init, intercept init, ...]) Fit linear model with Stochastic Gradient Descent.

```


```
# predict(X) Predict class labels for samples in X.

#-----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/geometric-intuition-1/
#-----
clf = SGDClassifier(class_weight='balanced', alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=42)
predict_and_plot_confusion_matrix(train_x_onehotCoding, train_y, cv_x_onehotCoding, cv_y, clf)
```


Log loss : 0.9554926080094617

Number of mis-classified points : 0.34022556390977443


----- Confusion matrix -----




----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----





#### 4.3.1.3. Feature Importance

In [66]:

```
def get_imp_feature_names(text, indices, removed_ind = []):
    word_present = 0
    tabulte_list = []
    incresingorder_ind = 0
    for i in indices:
        if i < train_gene_feature_onehotCoding.shape[1]:
            tabulte_list.append([incresingorder_ind, "Gene", "Yes"])
        elif i < 18:
            tabulte_list.append([incresingorder_ind, "Variation", "Yes"])
        if ((i > 17) & (i not in removed_ind)) :
            word = train_text_features[i]
            yes_no = True if word in text.split() else False
            if yes_no:
                word_present += 1
            tabulte_list.append([incresingorder_ind, train_text_features[i], yes_no])
            incresingorder_ind += 1
    print(word_present, "most important features are present in our query point")
    print("-"*50)
    print("The features that are most important of the ", predicted_cls[0], " class:")
    print(tabulate(tabulte_list, headers=["Index", "Feature name", "Present or Not"]))
```

##### 4.3.1.3.1. Correctly Classified point

In [67]:

```
# from tabulate import tabulate
clf = SGDClassifier(class_weight='balanced', alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=42)
clf.fit(train_x_onehotCoding, train_y)
test_point_index = 1
no_feature = 500
predicted_cls = sig_clf.predict(test_x_onehotCoding[test_point_index])
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:",
      np.round(sig_clf.predict_proba(test_x_onehotCoding[test_point_index]), 4))
print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-clf.coef_) [predicted_cls-1][:, :no_feature]
print("-"*50)
get_impfeature_names(indices[0],
test_df['TEXT'].iloc[test_point_index], test_df['Gene'].iloc[test_point_index], test_df['Variation'].iloc[test_point_index], no_feature)
```

```
Predicted Class : 4
Predicted Class Probabilities: [[3.766e-01 5.990e-02 5.900e-03 4.614e-01 2.100e-02 4.120e-02 3.120e-02
2.400e-03 4.000e-04]]
Actual Class : 3
-----
184 Text feature [1673] present in test data point [True]
481 Text feature [11] present in test data point [True]
Out of the top 500 features 2 are present in query point
```

##### 4.3.1.3.2. Incorrectly Classified point

In [68]:

```
test_point_index = 100
no_feature = 500
```

```

predicted_cls = sig_clf.predict(test_x_onehotCoding[test_point_index])
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:",
np.round(sig_clf.predict_proba(test_x_onehotCoding[test_point_index]),4))
print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-clf.coef_) [predicted_cls-1][:, :no_feature]
print("-"*50)
get_impfeature_names(indices[0],
test_df['TEXT'].iloc[test_point_index],test_df['Gene'].iloc[test_point_index],test_df['Variation']
.iloc[test_point_index], no_feature)

```

```

Predicted Class : 1
Predicted Class Probabilities: [[5.295e-01 5.900e-03 8.000e-04 4.429e-01 7.800e-03 8.700e-03 2.500
e-03
1.600e-03 3.000e-04]]
Actual Class : 4
-----
221 Text feature [11] present in test data point [True]
234 Text feature [13] present in test data point [True]
328 Text feature [125] present in test data point [True]
432 Text feature [163] present in test data point [True]
Out of the top 500 features 4 are present in query point

```

## 4.3.2. Without Class balancing

### 4.3.2.1. Hyper paramter tuning

In [69]:

```

# read more about SGDClassifier() at http://scikit-
learn.org/stable/modules/generated/sklearn.linear_model.SGDClassifier.html
# -----
# default parameters
# SGDClassifier(loss='hinge', penalty='l2', alpha=0.0001, l1_ratio=0.15, fit_intercept=True, max_i
ter=None, tol=None,
# shuffle=True, verbose=0, epsilon=0.1, n_jobs=1, random_state=None, learning_rate='optimal', eta0
=0.0, power_t=0.5,
# class_weight=None, warm_start=False, average=False, n_iter=None)

# some of methods
# fit(X, y[, coef_init, intercept_init, ...]) Fit linear model with Stochastic Gradient Descent.
# predict(X) Predict class labels for samples in X.

#-----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/geometric-in
tuition-1/
#-----

# find more about CalibratedClassifierCV here at http://scikit-
learn.org/stable/modules/generated/sklearn.calibration.CalibratedClassifierCV.html
# -----
# default paramters
# sklearn.calibration.CalibratedClassifierCV(base_estimator=None, method='sigmoid', cv=3)
#
# some of the methods of CalibratedClassifierCV()
# fit(X, y[, sample_weight]) Fit the calibrated model
# get_params([deep]) Get parameters for this estimator.
# predict(X) Predict the target of new samples.
# predict_proba(X) Posterior probabilities of classification
#-----
# video link:
#-----

alpha = [10 ** x for x in range(-6, 1)]
cv_log_error_array = []
for i in alpha:
    print("for alpha =", i)
    clf = SGDClassifier(alpha=i, penalty='l2', loss='log', random_state=42)
    clf.fit(train_x_onehotCoding, train_y)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_x_onehotCoding, train_y)

```

```

sig_clf_probs = sig_clf.predict_proba(cv_x_onehotCoding)
cv_log_error_array.append(log_loss(cv_y, sig_clf_probs, labels=clf.classes_, eps=1e-15))
print("Log Loss :", log_loss(cv_y, sig_clf_probs))

```

```

fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array, c='g')
for i, txt in enumerate(np.round(cv_log_error_array, 3)):
    ax.annotate((alpha[i], str(txt)), (alpha[i], cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

```

```

best_alpha = np.argmin(cv_log_error_array)
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=42)
clf.fit(train_x_onehotCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_onehotCoding, train_y)


predict_y = sig_clf.predict_proba(train_x_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:", log_loss(y_train,
predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(cv_x_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The cross validation log loss is:", log_lo
ss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(test_x_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:", log_loss(y_test, p
redict_y, labels=clf.classes_, eps=1e-15))

```

```

for alpha = 1e-06
Log Loss : 1.0416137856317276
for alpha = 1e-05
Log Loss : 1.0266312464450722
for alpha = 0.0001
Log Loss : 0.9961974290273851
for alpha = 0.001
Log Loss : 1.0756890922790818
for alpha = 0.01
Log Loss : 1.3733161040645743
for alpha = 0.1
Log Loss : 1.6475765253978356
for alpha = 1
Log Loss : 1.7816740956313173

```



```

For values of best alpha = 0.0001 The train log loss is: 0.4145289103332108
For values of best alpha = 0.0001 The cross validation log loss is: 0.9961974290273851
For values of best alpha = 0.0001 The test log loss is: 1.089069587425559

```

#### 4.3.2.2. Testing model with best hyper parameters

In [70]:

```

# read more about SGDClassifier() at http://scikit-learn.org/stable/modules/generated/sklearn.linear\_model.SGDClassifier.html

```

```
# -----
# default parameters
# SGDClassifier(loss='hinge', penalty='l2', alpha=0.0001, l1_ratio=0.15, fit_intercept=True, max_iter=None, tol=None,
# shuffle=True, verbose=0, epsilon=0.1, n_jobs=1, random_state=None, learning_rate='optimal', eta0=0.0, power_t=0.5,
# class_weight=None, warm_start=False, average=False, n_iter=None)

# some of methods
# fit(X, y[, coef_init, intercept_init, ...]) Fit linear model with Stochastic Gradient Descent.
# predict(X) Predict class labels for samples in X.


#-----
# video link:
#-----

clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=42)
predict_and_plot_confusion_matrix(train_x_onehotCoding, train_y, cv_x_onehotCoding, cv_y, clf)
```


Log loss : 0.9961974290273851

Number of mis-classified points : 0.31954887218045114


----- Confusion matrix -----




----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----







#### 4.3.2.3. Feature Importance, Correctly Classified point

In [71]:

```
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='log', random_state=42)
clf.fit(train_x_onehotCoding,train_y)
test_point_index = 1
no_feature = 500
predicted_cls = sig_clf.predict(test_x_onehotCoding[test_point_index])
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:",
np.round(sig_clf.predict_proba(test_x_onehotCoding[test_point_index]),4))
print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-clf.coef_) [predicted_cls-1][:,no_feature]
print("-"*50)
get_impfeature_names(indices[0],
test_df['TEXT'].iloc[test_point_index],test_df['Gene'].iloc[test_point_index],test_df['Variation']
.iloc[test_point_index], no_feature)
```

```
Predicted Class : 1
Predicted Class Probabilities: [[5.061e-01 2.340e-02 4.200e-03 3.611e-01 1.200e-02 2.990e-02 6.010
e-02
3.000e-03 2.000e-04]]
Actual Class : 3
-----
216 Text feature [13] present in test data point [True]
290 Text feature [06] present in test data point [True]
443 Text feature [1008] present in test data point [True]
Out of the top 500 features 3 are present in query point
```

#### 4.3.2.4. Feature Importance, Inorrectly Classified point

In [72]:

```
test_point_index = 100
no_feature = 500
predicted_cls = sig_clf.predict(test_x_onehotCoding[test_point_index])
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:",
np.round(sig_clf.predict_proba(test_x_onehotCoding[test_point_index]),4))
print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-clf.coef_) [predicted_cls-1][:,no_feature]
print("-"*50)
get_impfeature_names(indices[0],
test_df['TEXT'].iloc[test_point_index],test_df['Gene'].iloc[test_point_index],test_df['Variation']
.iloc[test_point_index], no_feature)
```

```
Predicted Class : 4
Predicted Class Probabilities: [[4.256e-01 1.590e-02 2.000e-04 5.396e-01 7.900e-03 5.800e-03 2.900
e-03
2.000e-03 1.000e-04]]
Actual Class : 4
-----
357 Text feature [125] present in test data point [True]
359 Text feature [1met] present in test data point [True]
```

```
387 Text feature [11] present in test data point [True]
425 Text feature [04] present in test data point [True]
Out of the top 500 features 4 are present in query point
```

## 4.4. Linear Support Vector Machines

### 4.4.1. Hyper paramter tuning

In [73]:

```
# read more about support vector machines with linear kernals here http://scikit-learn.org/stable/modules/generated/sklearn.svm.SVC.html

# -----
# default parameters
# SVC(C=1.0, kernel='rbf', degree=3, gamma='auto', coef0=0.0, shrinking=True, probability=False, tol=0.001,
# cache_size=200, class_weight=None, verbose=False, max_iter=-1, decision_function_shape='ovr', random_state=None)

# Some of methods of SVM()
# fit(X, y, [sample_weight]) Fit the SVM model according to the given training data.
# predict(X) Perform classification on samples in X.
# -----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/mathematical-derivation-copy-8/
# -----

# find more about CalibratedClassifierCV here at http://scikit-learn.org/stable/modules/generated/sklearn.calibration.CalibratedClassifierCV.html
# -----
# default paramters
# sklearn.calibration.CalibratedClassifierCV(base_estimator=None, method='sigmoid', cv=3)
#
# some of the methods of CalibratedClassifierCV()
# fit(X, y[, sample_weight]) Fit the calibrated model
# get_params([deep]) Get parameters for this estimator.
# predict(X) Predict the target of new samples.
# predict_proba(X) Posterior probabilities of classification
#-----
# video link:
#-----

alpha = [10 ** x for x in range(-5, 3)]
cv_log_error_array = []
for i in alpha:
    print("for C =", i)
    # clf = SVC(C=i, kernel='linear', probability=True, class_weight='balanced')
    clf = SGDClassifier(class_weight='balanced', alpha=i, penalty='l2', loss='hinge', random_state=42)
    clf.fit(train_x_onehotCoding, train_y)
    sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
    sig_clf.fit(train_x_onehotCoding, train_y)
    sig_clf_probs = sig_clf.predict_proba(cv_x_onehotCoding)
    cv_log_error_array.append(log_loss(cv_y, sig_clf_probs, labels=clf.classes_, eps=1e-15))
    print("Log Loss :", log_loss(cv_y, sig_clf_probs))

fig, ax = plt.subplots()
ax.plot(alpha, cv_log_error_array, c='g')
for i, txt in enumerate(np.round(cv_log_error_array, 3)):
    ax.annotate((alpha[i], str(txt)), (alpha[i], cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()

best_alpha = np.argmin(cv_log_error_array)
# clf = SVC(C=i, kernel='linear', probability=True, class_weight='balanced')
clf = SGDClassifier(class_weight='balanced', alpha=alpha[best_alpha], penalty='l2', loss='hinge', r
```

```

andom_state=42)
clf.fit(train_x_onehotCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_onehotCoding, train_y)


predict_y = sig_clf.predict_proba(train_x_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The train log loss is:", log_loss(y_train,
predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(cv_x_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The cross validation log loss is:", log_lo
ss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(test_x_onehotCoding)
print('For values of best alpha = ', alpha[best_alpha], "The test log loss is:", log_loss(y_test, p
redict_y, labels=clf.classes_, eps=1e-15))

```

```

for C = 1e-05
Log Loss : 1.0277883889967174
for C = 0.0001
Log Loss : 1.017711976873179
for C = 0.001
Log Loss : 0.98672780417483
for C = 0.01
Log Loss : 1.1863350195666134
for C = 0.1
Log Loss : 1.6614809870077587
for C = 1
Log Loss : 1.81129299984073
for C = 10
Log Loss : 1.8112857229231774
for C = 100
Log Loss : 1.8112887384029561

```



```

For values of best alpha = 0.001 The train log loss is: 0.5486032324500404
For values of best alpha = 0.001 The cross validation log loss is: 0.98672780417483
For values of best alpha = 0.001 The test log loss is: 1.1185141557394398

```

#### 4.4.2. Testing model with best hyper parameters

In [74]:

```

# read more about support vector machines with linear kernels here http://scikit-
learn.org/stable/modules/generated/sklearn.svm.SVC.html

# -----
# default parameters
# SVC(C=1.0, kernel='rbf', degree=3, gamma='auto', coef0=0.0, shrinking=True, probability=False, t
ol=0.001,
# cache_size=200, class_weight=None, verbose=False, max_iter=-1, decision_function_shape='ovr', ra
ndom_state=None)

# Some of methods of SVM()
# fit(X, y, [sample_weight]) Fit the SVM model according to the given training data.
# predict(X) Perform classification on samples in X.
# -----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-
online/lessons/mathematical-derivations-part-8/

```


# -----

```
# clf = SVC(C=alpha[best_alpha],kernel='linear',probability=True, class_weight='balanced')
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='hinge',
random_state=42,class_weight='balanced')
predict_and_plot_confusion_matrix(train_x_onehotCoding, train_y,cv_x_onehotCoding,cv_y, clf)
```


Log loss : 0.98672780417483

Number of mis-classified points : 0.34398496240601506


----- Confusion matrix -----




----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----





### 4.3.3. Feature Importance

#### 4.3.3.1. For Correctly classified point

In [75]:

```
clf = SGDClassifier(alpha=alpha[best_alpha], penalty='l2', loss='hinge', random_state=42)
clf.fit(train_x_onehotCoding, train_y)
test_point_index = 1
# test_point_index = 100
no_feature = 500
predicted_cls = sig_clf.predict(test_x_onehotCoding[test_point_index])
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:",
      np.round(sig_clf.predict_proba(test_x_onehotCoding[test_point_index]), 4))
print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-clf.coef_)[predicted_cls-1][:, :no_feature]
print("-"*50)
get_impfeature_names(indices[0],
test_df['TEXT'].iloc[test_point_index], test_df['Gene'].iloc[test_point_index], test_df['Variation']
.iloc[test_point_index], no_feature)
```

```
Predicted Class : 1
Predicted Class Probabilities: [[0.4849 0.0321 0.0094 0.3564 0.012 0.0293 0.0739 0.001 0.001 ]]
Actual Class : 3
-----
300 Text feature [06] present in test data point [True]
391 Text feature [18] present in test data point [True]
465 Text feature [13] present in test data point [True]
468 Text feature [1008] present in test data point [True]
Out of the top 500 features 4 are present in query point
```

#### 4.3.3.2. For Incorrectly classified point

In [76]:

```
test_point_index = 100
no_feature = 500
predicted_cls = sig_clf.predict(test_x_onehotCoding[test_point_index])
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:",
      np.round(sig_clf.predict_proba(test_x_onehotCoding[test_point_index]), 4))
print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-clf.coef_)[predicted_cls-1][:, :no_feature]
print("-"*50)
get_impfeature_names(indices[0],
test_df['TEXT'].iloc[test_point_index], test_df['Gene'].iloc[test_point_index], test_df['Variation']
.iloc[test_point_index], no_feature)
```

```
Predicted Class : 1
Predicted Class Probabilities: [[0.4911 0.0529 0.0022 0.3522 0.0442 0.0148 0.0378 0.0024 0.0023]]
Actual Class : 4
-----
391 Text feature [18] present in test data point [True]
419 Text feature [158] present in test data point [True]
465 Text feature [13] present in test data point [True]
Out of the top 500 features 3 are present in query point
```

## 4.5 Random Forest Classifier

### 4.5.1 Hyperparameter Tuning with Grid Search

### 4.5.1. Hyper paramter tuning (With One hot Encoding)

In [77]:

```
# -----
# default parameters
# sklearn.ensemble.RandomForestClassifier(n_estimators=10, criterion='gini', max_depth=None, min_
amples_split=2,
# min_samples_leaf=1, min_weight_fraction_leaf=0.0, max_features='auto', max_leaf_nodes=None, min_
impurity_decrease=0.0,
# min_impurity_split=None, bootstrap=True, oob_score=False, n_jobs=1, random_state=None,
verbose=0, warm_start=False,
# class_weight=None)

# Some of methods of RandomForestClassifier()
# fit(X, y, [sample_weight]) Fit the SVM model according to the given training data.
# predict(X) Perform classification on samples in X.
# predict_proba (X) Perform classification on samples in X.

# some of attributes of RandomForestClassifier()
# feature_importances_ : array of shape = [n_features]
# The feature importances (the higher, the more important the feature).

# -----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/random-forest-and-their-construction-2/
# -----

# find more about CalibratedClassifierCV here at http://scikit-learn.org/stable/modules/generated/sklearn.calibration.CalibratedClassifierCV.html
# -----
# default paramters
# sklearn.calibration.CalibratedClassifierCV(base_estimator=None, method='sigmoid', cv=3)
#
# some of the methods of CalibratedClassifierCV()
# fit(X, y[, sample_weight]) Fit the calibrated model
# get_params([deep]) Get parameters for this estimator.
# predict(X) Predict the target of new samples.
# predict_proba(X) Posterior probabilities of classification
#-----
# video link:
#-----

alpha = [100,200,500,1000,2000]
max_depth = [5, 10]
cv_log_error_array = []
for i in alpha:
    for j in max_depth:
        print("for n_estimators =", i,"and max depth = ", j)
        clf = RandomForestClassifier(n_estimators=i, criterion='gini', max_depth=j, random_state=42
, n_jobs=-1)
        clf.fit(train_x_onehotCoding, train_y)
        sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
        sig_clf.fit(train_x_onehotCoding, train_y)
        sig_clf_probs = sig_clf.predict_proba(cv_x_onehotCoding)
        cv_log_error_array.append(log_loss(cv_y, sig_clf_probs, labels=clf.classes_, eps=1e-15))
        print("Log Loss :",log_loss(cv_y, sig_clf_probs))

'''fig, ax = plt.subplots()
features = np.dot(np.array(alpha)[:,None],np.array(max_depth)[None]).ravel()
ax.plot(features, cv_log_error_array,c='g')
for i, txt in enumerate(np.round(cv_log_error_array,3)):
    ax.annotate((alpha[int(i/2)],max_depth[int(i%2)],str(txt)),
(features[i],cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()
'''

best_alpha = np.argmin(cv_log_error_array)
clf = RandomForestClassifier(n_estimators=alpha[int(best_alpha/2)], criterion='gini', max_depth=max_
depth[int(best_alpha%2)], random_state=42, n_jobs=-1)
clf.fit(train_x_onehotCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
```

```

sig_clf.fit(train_x_onehotCoding, train_y)

predict_y = sig_clf.predict_proba(train_x_onehotCoding)
print('For values of best estimator = ', alpha[int(best_alpha/2)], "The train log loss
is:", log_loss(y_train, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(cv_x_onehotCoding)
print('For values of best estimator = ', alpha[int(best_alpha/2)], "The cross validation log loss
is:", log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(test_x_onehotCoding)
print('For values of best estimator = ', alpha[int(best_alpha/2)], "The test log loss
is:", log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))

for n_estimators = 100 and max depth = 5
Log Loss : 1.1871913034149577
for n_estimators = 100 and max depth = 10
Log Loss : 1.2008652546844876
for n_estimators = 200 and max depth = 5
Log Loss : 1.177984271330672
for n_estimators = 200 and max depth = 10
Log Loss : 1.1857118498130315
for n_estimators = 500 and max depth = 5
Log Loss : 1.1603839257641497
for n_estimators = 500 and max depth = 10
Log Loss : 1.1797182123190368
for n_estimators = 1000 and max depth = 5
Log Loss : 1.1573728839303261
for n_estimators = 1000 and max depth = 10
Log Loss : 1.1715823371359846
for n_estimators = 2000 and max depth = 5
Log Loss : 1.1548718022616935
for n_estimators = 2000 and max depth = 10
Log Loss : 1.1730732218915103
For values of best estimator = 2000 The train log loss is: 0.8545244378444445
For values of best estimator = 2000 The cross validation log loss is: 1.1548718014820822
For values of best estimator = 2000 The test log loss is: 1.205521549726899

```

#### 4.5.2. Testing model with best hyper parameters (One Hot Encoding)

In [78]:

```

# -----
# default parameters
# sklearn.ensemble.RandomForestClassifier(n_estimators=10, criterion='gini', max_depth=None, min_s
amples_split=2,
# min_samples_leaf=1, min_weight_fraction_leaf=0.0, max_features='auto', max_leaf_nodes=None, min_
impurity_decrease=0.0,
# min_impurity_split=None, bootstrap=True, oob_score=False, n_jobs=1, random_state=None,
verbose=0, warm_start=False,
# class_weight=None)

# Some of methods of RandomForestClassifier()
# fit(X, y, [sample_weight]) Fit the SVM model according to the given training data.
# predict(X) Perform classification on samples in X.
# predict_proba (X) Perform classification on samples in X.

# some of attributes of RandomForestClassifier()
# feature_importances_ : array of shape = [n_features]
# The feature importances (the higher, the more important the feature).

# -----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/random-forest-and-their-construction-2/
# -----


clf = RandomForestClassifier(n_estimators=alpha[int(best_alpha/2)], criterion='gini', max_depth=max_
depth[int(best_alpha/2)], random_state=42, n_jobs=-1)
predict_and_plot_confusion_matrix(train_x_onehotCoding, train_y, cv_x_onehotCoding, cv_y, clf)

```


```

Log loss : 1.1548718022616937
Number of mis-classified points : 0.40977443609022557
----- Confusion matrix -----


```



----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----



## 4.5.3. Feature Importance

### 4.5.3.1. Correctly Classified point

In [79]:



```
# test_point_index = 10
clf = RandomForestClassifier(n_estimators=alpha[int(best_alpha/2)], criterion='gini', max_depth=max_depth[int(best_alpha%2)], random_state=42, n_jobs=-1)
clf.fit(train_x_onehotCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_onehotCoding, train_y)

test_point_index = 1
no_feature = 100
predicted_cls = sig_clf.predict(test_x_onehotCoding[test_point_index])
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:",
np.round(sig_clf.predict_proba(test_x_onehotCoding[test_point_index]),4))
print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-clf.feature_importances_)
print("--*50)
get_impfeature_names(indices[:no_feature], test_df['TEXT'].iloc[test_point_index],test_df['Gene'].iloc[test_point_index],test_df['Variation'].iloc[test_point_index], no_feature)
```

```
Predicted Class : 1
Predicted Class Probabilities: [[0.4156 0.061 0.0179 0.2158 0.092 0.0905 0.099 0.0065 0.0018]]
Actual Class : 3
-----
4 Text feature [1g] present in test data point [True]
26 Text feature [141] present in test data point [True]
35 Text feature [117] present in test data point [True]
71 Text feature [17] present in test data point [True]
92 Text feature [12] present in test data point [True]
96 Text feature [1673] present in test data point [True]
Out of the top 100 features 6 are present in query point
```

#### 4.5.3.2. Inorrectly Classified point

In [80]:

```
test_point_index = 100
no_feature = 100
predicted_cls = sig_clf.predict(test_x_onehotCoding[test_point_index])
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:",
np.round(sig_clf.predict_proba(test_x_onehotCoding[test_point_index]),4))
print("Actuall Class :", test_y[test_point_index])
indices = np.argsort(-clf.feature_importances_)
print("--*50)
get_impfeature_names(indices[:no_feature], test_df['TEXT'].iloc[test_point_index],test_df['Gene'].iloc[test_point_index],test_df['Variation'].iloc[test_point_index], no_feature)
```

```
Predicted Class : 1
Predicted Class Probabilities: [[0.3892 0.0565 0.0106 0.3033 0.0595 0.0529 0.1149 0.0069 0.0062]]
Actuall Class : 4
-----
16 Text feature [1c] present in test data point [True]
26 Text feature [141] present in test data point [True]
71 Text feature [17] present in test data point [True]
76 Text feature [125] present in test data point [True]
92 Text feature [12] present in test data point [True]
Out of the top 100 features 5 are present in query point
```

#### 4.5.3. Hyper paramter tuning (With Response Coding)

In [81]:

```
# -----
# default parameters
# sklearn.ensemble.RandomForestClassifier(n_estimators=10, criterion='gini', max_depth=None, min_samples_split=2,
# min_samples_leaf=1, min_weight_fraction_leaf=0.0, max_features='auto', max_leaf_nodes=None, min_impurity_decrease=0.0,
# min_impurity_split=None, bootstrap=True, oob_score=False, n_jobs=1, random_state=None,
verbose=0, warm_start=False,
# class_weight=None)
```

```

# class_weight=None,

# Some of methods of RandomForestClassifier()
# fit(X, y, [sample_weight]) Fit the SVM model according to the given training data.
# predict(X) Perform classification on samples in X.
# predict_proba(X) Perform classification on samples in X.

# some of attributes of RandomForestClassifier()
# feature_importances_ : array of shape = [n_features]
# The feature importances (the higher, the more important the feature).

# -----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/random-forest-and-their-construction-2/
# -----

# find more about CalibratedClassifierCV here at http://scikit-learn.org/stable/modules/generated/sklearn.calibration.CalibratedClassifierCV.html
# -----
# default paramters
# sklearn.calibration.CalibratedClassifierCV(base_estimator=None, method='sigmoid', cv=3)
#
# some of the methods of CalibratedClassifierCV()
# fit(X, y[, sample_weight]) Fit the calibrated model
# get_params([deep]) Get parameters for this estimator.
# predict(X) Predict the target of new samples.
# predict_proba(X) Posterior probabilities of classification
#-----
# video link:
#-----

alpha = [10,50,100,200,500,1000]
max_depth = [2,3,5,10]
cv_log_error_array = []
for i in alpha:
    for j in max_depth:
        print("for n_estimators = ", i,"and max depth = ", j)
        clf = RandomForestClassifier(n_estimators=i, criterion='gini', max_depth=j, random_state=42, n_jobs=-1)
        clf.fit(train_x_responseCoding, train_y)
        sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
        sig_clf.fit(train_x_responseCoding, train_y)
        sig_clf_probs = sig_clf.predict_proba(cv_x_responseCoding)
        cv_log_error_array.append(log_loss(cv_y, sig_clf_probs, labels=clf.classes_, eps=1e-15))
        print("Log Loss :",log_loss(cv_y, sig_clf_probs))
'''
fig, ax = plt.subplots()
features = np.dot(np.array(alpha)[:,None],np.array(max_depth)[None]).ravel()
ax.plot(features, cv_log_error_array,c='g')
for i, txt in enumerate(np.round(cv_log_error_array,3)):
    ax.annotate((alpha[int(i/4)],max_depth[int(i%4)],str(txt)),
        (features[i],cv_log_error_array[i]))
plt.grid()
plt.title("Cross Validation Error for each alpha")
plt.xlabel("Alpha i's")
plt.ylabel("Error measure")
plt.show()
'''

best_alpha = np.argmin(cv_log_error_array)
clf = RandomForestClassifier(n_estimators=alpha[int(best_alpha/4)], criterion='gini', max_depth=max_depth[int(best_alpha%4)], random_state=42, n_jobs=-1)
clf.fit(train_x_responseCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_responseCoding, train_y)

predict_y = sig_clf.predict_proba(train_x_responseCoding)
print('For values of best alpha = ', alpha[int(best_alpha/4)], "The train log loss is:",log_loss(y_train, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(cv_x_responseCoding)
print('For values of best alpha = ', alpha[int(best_alpha/4)], "The cross validation log loss is:"
,log_loss(y_cv, predict_y, labels=clf.classes_, eps=1e-15))
predict_y = sig_clf.predict_proba(test_x_responseCoding)
print('For values of best alpha = ', alpha[int(best_alpha/4)], "The test log loss is:",log_loss(y_test, predict_y, labels=clf.classes_, eps=1e-15))

```

```

for n_estimators = 10 and max depth = 2
Log Loss : 2.139267104191326
for n_estimators = 10 and max depth = 3
Log Loss : 1.6299104374172773
for n_estimators = 10 and max depth = 5
Log Loss : 1.41424870577866
for n_estimators = 10 and max depth = 10
Log Loss : 1.6270134996125234
for n_estimators = 50 and max depth = 2
Log Loss : 1.7310221225998332
for n_estimators = 50 and max depth = 3
Log Loss : 1.4041832606738556
for n_estimators = 50 and max depth = 5
Log Loss : 1.3244255868561363
for n_estimators = 50 and max depth = 10
Log Loss : 1.7757140475003752
for n_estimators = 100 and max depth = 2
Log Loss : 1.5972056754778616
for n_estimators = 100 and max depth = 3
Log Loss : 1.4323699055270096
for n_estimators = 100 and max depth = 5
Log Loss : 1.2507739261153945
for n_estimators = 100 and max depth = 10
Log Loss : 1.726293381836022
for n_estimators = 200 and max depth = 2
Log Loss : 1.6403073593465964
for n_estimators = 200 and max depth = 3
Log Loss : 1.4598520984779337
for n_estimators = 200 and max depth = 5
Log Loss : 1.2971390642835883
for n_estimators = 200 and max depth = 10
Log Loss : 1.6868808157429092
for n_estimators = 500 and max depth = 2
Log Loss : 1.6665956169330705
for n_estimators = 500 and max depth = 3
Log Loss : 1.510857860851123
for n_estimators = 500 and max depth = 5
Log Loss : 1.3145229485201817
for n_estimators = 500 and max depth = 10
Log Loss : 1.713294369727822
for n_estimators = 1000 and max depth = 2
Log Loss : 1.6686845838021007
for n_estimators = 1000 and max depth = 3
Log Loss : 1.5198573589099125
for n_estimators = 1000 and max depth = 5
Log Loss : 1.304623884700354
for n_estimators = 1000 and max depth = 10
Log Loss : 1.709877150168174
For values of best alpha = 100 The train log loss is: 0.05480299322089486
For values of best alpha = 100 The cross validation log loss is: 1.250773926115395
For values of best alpha = 100 The test log loss is: 1.3004141198588886

```

#### 4.5.4. Testing model with best hyper parameters (Response Coding)

In [82]:

```

# -----
# default parameters
# sklearn.ensemble.RandomForestClassifier(n_estimators=10, criterion='gini', max_depth=None, min_
samples_split=2,
# min_samples_leaf=1, min_weight_fraction_leaf=0.0, max_features='auto', max_leaf_nodes=None, min_
impurity_decrease=0.0,
# min_impurity_split=None, bootstrap=True, oob_score=False, n_jobs=1, random_state=None,
verbose=0, warm_start=False,
# class_weight=None)

# Some of methods of RandomForestClassifier()
# fit(X, y, [sample_weight]) Fit the SVM model according to the given training data.
# predict(X) Perform classification on samples in X.
# predict_proba (X) Perform classification on samples in X.

# some of attributes of RandomForestClassifier()
# feature_importances_ : array of shape = [n_features]
# The feature importances (the higher, the more important the feature).

```


```
# -----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/random-forest-and-their-construction-2/
# -----

clf = RandomForestClassifier(max_depth=max_depth[int(best_alpha%4)],
n_estimators=alpha[int(best_alpha/4)], criterion='gini', max_features='auto',random_state=42)
predict_and_plot_confusion_matrix(train_x_responseCoding, train_y,cv_x_responseCoding,cv_y, clf)
```


Log loss : 1.250773926115395

Number of mis-classified points : 0.4473684210526316


----- Confusion matrix -----




----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----





## 4.5.5. Feature Importance

### 4.5.5.1. Correctly Classified point

In [83]:

```
clf = RandomForestClassifier(n_estimators=alpha[int(best_alpha/4)], criterion='gini', max_depth=max_
_depth[int(best_alpha%4)], random_state=42, n_jobs=-1)
clf.fit(train_x_responseCoding, train_y)
sig_clf = CalibratedClassifierCV(clf, method="sigmoid")
sig_clf.fit(train_x_responseCoding, train_y)
```

```
test_point_index = 1
no_feature = 27
predicted_cls = sig_clf.predict(test_x_responseCoding[test_point_index].reshape(1,-1))
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:",
np.round(sig_clf.predict_proba(test_x_responseCoding[test_point_index].reshape(1,-1)),4))
print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-clf.feature_importances_)
print("-"*50)
for i in indices:
    if i<9:
        print("Gene is important feature")
    elif i<18:
        print("Variation is important feature")
    else:
        print("Text is important feature")
```

Predicted Class : 4

Predicted Class Probabilities: [[0.2071 0.0236 0.1277 0.4264 0.065 0.0763 0.0118 0.0275 0.0346]]

Actual Class : 3

```
-----
Variation is important feature
Variation is important feature
Variation is important feature
Variation is important feature
Gene is important feature
Variation is important feature
Variation is important feature
Text is important feature
Text is important feature
Text is important feature
Gene is important feature
Gene is important feature
Text is important feature
Text is important feature
Gene is important feature
Variation is important feature
Text is important feature
Gene is important feature
Variation is important feature
Variation is important feature
Text is important feature
Gene is important feature
Text is important feature
Text is important feature
Gene is important feature
Gene is important feature
Gene is important feature
```

### 4.5.5.2. Incorrectly Classified point

In [84]:

```
test_point_index = 100
predicted_cls = sig_clf.predict(test_x_responseCoding[test_point_index].reshape(1,-1))
print("Predicted Class :", predicted_cls[0])
print("Predicted Class Probabilities:",
np.round(sig_clf.predict_proba(test_x_responseCoding[test_point_index].reshape(1,-1)),4))
print("Actual Class :", test_y[test_point_index])
indices = np.argsort(-clf.feature_importances_)
print("--*50)
for i in indices:
    if i<9:
        print("Gene is important feature")
    elif i<18:
        print("Variation is important feature")
    else:
        print("Text is important feature")
```

Predicted Class : 4

Predicted Class Probabilities: [[0.2788 0.021 0.1708 0.3726 0.0351 0.0661 0.0096 0.0172 0.0289]]

Actual Class : 4

-----

Variation is important feature  
Variation is important feature  
Variation is important feature  
Variation is important feature  
Gene is important feature  
Variation is important feature  
Variation is important feature  
Text is important feature  
Text is important feature  
Text is important feature  
Gene is important feature  
Gene is important feature  
Text is important feature  
Text is important feature  
Gene is important feature  
Variation is important feature  
Text is important feature  
Gene is important feature  
Variation is important feature  
Variation is important feature  
Text is important feature  
Gene is important feature  
Text is important feature  
Text is important feature  
Gene is important feature  
Gene is important feature  
Gene is important feature

## 4.7 Stack the models

### 4.7.1 testing with hyper parameter tuning

In [85]:

```
# read more about SGDClassifier() at http://scikit-
learn.org/stable/modules/generated/sklearn.linear_model.SGDClassifier.html
# -----
# default parameters
# SGDClassifier(loss='hinge', penalty='l2', alpha=0.0001, l1_ratio=0.15, fit_intercept=True, max_i
ter=None, tol=None,
# shuffle=True, verbose=0, epsilon=0.1, n_jobs=1, random_state=None, learning_rate='optimal', eta0
=0.0, power_t=0.5,
# class_weight=None, warm_start=False, average=False, n_iter=None)

# some of methods
# fit(X, y[, coef_init, intercept_init, ...]) Fit linear model with Stochastic Gradient Descent.
# predict(X) Predict class labels for samples in X.

#-----
# video link: https://www.appliedaicaourse.com/course/applied-ai-course-online/lessons/geometric-in
```

```

tuition-1/
#-----

# read more about support vector machines with linear kernals here http://scikit-learn.org/stable/modules/generated/sklearn.svm.SVC.html
# -----
# default parameters
# SVC(C=1.0, kernel='rbf', degree=3, gamma='auto', coef0=0.0, shrinking=True, probability=False, tol=0.001,
# cache_size=200, class_weight=None, verbose=False, max_iter=-1, decision_function_shape='ovr', random_state=None)

# Some of methods of SVM()
# fit(X, y, [sample_weight]) Fit the SVM model according to the given training data.
# predict(X) Perform classification on samples in X.
# -----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/mathematical-derivation-copy-8/
# -----

# read more about support vector machines with linear kernals here http://scikit-learn.org/stable/modules/generated/sklearn.ensemble.RandomForestClassifier.html
# -----
# default parameters
# sklearn.ensemble.RandomForestClassifier(n_estimators=10, criterion='gini', max_depth=None, min_samples_split=2,
# min_samples_leaf=1, min_weight_fraction_leaf=0.0, max_features='auto', max_leaf_nodes=None, min_impurity_decrease=0.0,
# min_impurity_split=None, bootstrap=True, oob_score=False, n_jobs=1, random_state=None, verbose=0, warm_start=False,
# class_weight=None)

# Some of methods of RandomForestClassifier()
# fit(X, y, [sample_weight]) Fit the SVM model according to the given training data.
# predict(X) Perform classification on samples in X.
# predict_proba(X) Perform classification on samples in X.

# some of attributes of RandomForestClassifier()
# feature_importances_ : array of shape = [n_features]
# The feature importances (the higher, the more important the feature).

# -----
# video link: https://www.appliedaicourse.com/course/applied-ai-course-online/lessons/random-forest-and-their-construction-2/
# -----

clf1 = SGDClassifier(alpha=0.001, penalty='l2', loss='log', class_weight='balanced', random_state=0)
clf1.fit(train_x_onehotCoding, train_y)
sig_clf1 = CalibratedClassifierCV(clf1, method="sigmoid")

clf2 = SGDClassifier(alpha=1, penalty='l2', loss='hinge', class_weight='balanced', random_state=0)
clf2.fit(train_x_onehotCoding, train_y)
sig_clf2 = CalibratedClassifierCV(clf2, method="sigmoid")

clf3 = MultinomialNB(alpha=0.001)
clf3.fit(train_x_onehotCoding, train_y)
sig_clf3 = CalibratedClassifierCV(clf3, method="sigmoid")

sig_clf1.fit(train_x_onehotCoding, train_y)
print("Logistic Regression : Log Loss: %0.2f" % (log_loss(cv_y, sig_clf1.predict_proba(cv_x_onehotCoding))))
sig_clf2.fit(train_x_onehotCoding, train_y)
print("Support vector machines : Log Loss: %0.2f" % (log_loss(cv_y, sig_clf2.predict_proba(cv_x_onehotCoding))))
sig_clf3.fit(train_x_onehotCoding, train_y)
print("Naive Bayes : Log Loss: %0.2f" % (log_loss(cv_y, sig_clf3.predict_proba(cv_x_onehotCoding))))
print("-"*50)
alpha = [0.0001, 0.001, 0.01, 0.1, 1, 10]
best_alpha = 999
for i in alpha:
    lr = LogisticRegression(C=i)

```

```

sclf = StackingClassifier(classifiers=[sig_clf1, sig_clf2, sig_clf3], meta_classifier=lr, use_p
robas=True)
sclf.fit(train_x_onehotCoding, train_y)
print("Stacking Classifier : for the value of alpha: %f Log Loss: %0.3f" % (i, log_loss(cv_y, sc
lf.predict_proba(cv_x_onehotCoding))))
log_error = log_loss(cv_y, sclf.predict_proba(cv_x_onehotCoding))
if best_alpha > log_error:
    best_alpha = log_error

```

Logistic Regression : Log Loss: 0.95  
Support vector machines : Log Loss: 1.81  
Naive Bayes : Log Loss: 1.15

-----  
Stacking Classifier : for the value of alpha: 0.000100 Log Loss: 2.177  
Stacking Classifier : for the value of alpha: 0.001000 Log Loss: 2.027  
Stacking Classifier : for the value of alpha: 0.010000 Log Loss: 1.474  
Stacking Classifier : for the value of alpha: 0.100000 Log Loss: 1.104  
Stacking Classifier : for the value of alpha: 1.000000 Log Loss: 1.268  
Stacking Classifier : for the value of alpha: 10.000000 Log Loss: 1.664

## 4.7.2 testing the model with the best hyper parameters

In [86]:

```

lr = LogisticRegression(C=0.1)
sclf = StackingClassifier(classifiers=[sig_clf1, sig_clf2, sig_clf3], meta_classifier=lr, use_proba
s=True)
sclf.fit(train_x_onehotCoding, train_y)

log_error = log_loss(train_y, sclf.predict_proba(train_x_onehotCoding))
print("Log loss (train) on the stacking classifier :", log_error)

log_error = log_loss(cv_y, sclf.predict_proba(cv_x_onehotCoding))
print("Log loss (CV) on the stacking classifier :", log_error)


log_error = log_loss(test_y, sclf.predict_proba(test_x_onehotCoding))
print("Log loss (test) on the stacking classifier :", log_error)

print("Number of missclassified point :", np.count_nonzero((sclf.predict(test_x_onehotCoding)-
test_y))/test_y.shape[0])
plot_confusion_matrix(test_y=test_y, predict_y=sclf.predict(test_x_onehotCoding))

```

Log loss (train) on the stacking classifier : 0.5919636909006337  
Log loss (CV) on the stacking classifier : 1.103928043162243  
Log loss (test) on the stacking classifier : 1.1820319422217422  
Number of missclassified point : 0.3819548872180451

----- Confusion matrix -----




----- Precision matrix (Column Sum=1) -----





----- Recall matrix (Row sum=1) -----



### 4.7.3 Maximum Voting classifier

In [87]:

```
#Refer:http://scikit-learn.org/stable/modules/generated/sklearn.ensemble.VotingClassifier.html
from sklearn.ensemble import VotingClassifier
vclf = VotingClassifier(estimators=[('lr', sig_clf1), ('svc', sig_clf2), ('rf', sig_clf3)], voting='soft')
vclf.fit(train_x_onehotCoding, train_y)
print("Log loss (train) on the VotingClassifier :", log_loss(train_y,
vclf.predict_proba(train_x_onehotCoding)))
print("Log loss (CV) on the VotingClassifier :", log_loss(cv_y,
vclf.predict_proba(cv_x_onehotCoding)))
print("Log loss (test) on the VotingClassifier :", log_loss(test_y,
vclf.predict_proba(test_x_onehotCoding)))
print("Number of missclassified point :", np.count_nonzero((vclf.predict(test_x_onehotCoding)-
test_y))/test_y.shape[0])
plot_confusion_matrix(test_y=test_y, predict_y=vclf.predict(test_x_onehotCoding))
```


Log loss (train) on the VotingClassifier : 0.8469761018953182

Log loss (CV) on the VotingClassifier : 1.1281406017858757

Log loss (test) on the VotingClassifier : 1.2041189385686248


Number of missclassified point : 0.3774436090225564

----- Confusion matrix -----






----- Precision matrix (Column Sum=1) -----



----- Recall matrix (Row sum=1) -----



## 5. Conclusion

1. Applied TfidfVectorizer on all three features ; Gene, Variation and text and taken top 2000 words with ngram\_range=(1,4)

In [89]:

```
from prettytable import PrettyTable
```

```

x = PrettyTable()

x.field_names = ["Vectorization", "Model", "Train Loss", "CV Loss", "Test Loss", "Percentage Misclassified"]

x.add_row(["OneHotEnCoding", "NaiveBayes", 0.58, 1.15, 1.23, 37.59])
x.add_row(["ResponseCoding", "KNN", 0.66, 1.01, 1.08, 35.33])
x.add_row(["OneHotEnCoding_ClassBalancing", "LogisticRegression", 0.67, 0.95, 1.08, 34.02])
x.add_row(["OneHotEnCoding_Without_ClassBalancing", "LogisticRegression", 0.41, 0.99, 1.08, 31.95])
x.add_row(["OneHotEnCoding", "LinearSVM", 0.57, 0.98, 1.11, 34.39])
x.add_row(["OneHotEnCoding", "RandomForest", 0.85, 1.15, 1.20, 40.97])
x.add_row(["ResponseCoding", "RandomForest", 0.05, 1.25, 1.30, 44.73])
x.add_row(["OneHotEnCoding", "Stacking", 0.59, 1.10, 1.18, 38.19])
x.add_row(["OneHotEnCoding", "Voting", 0.84, 1.12, 1.20, 37.74])

print(x)

```

Vectorization	Model	Train Loss	CV Loss	Test Loss	Percentage Misclassified
OneHotEnCoding	NaiveBayes	0.58	1.15	1.23	37.59
ResponseCoding	KNN	0.66	1.01	1.08	35.33
OneHotEnCoding_ClassBalancing	LogisticRegression	0.67	0.95	1.08	34.02
OneHotEnCoding_Without_ClassBalancing	LogisticRegression	0.41	0.99	1.08	31.95
OneHotEnCoding	LinearSVM	0.57	0.98	1.11	34.39
OneHotEnCoding	RandomForest	0.85	1.15	1.2	40.97
ResponseCoding	RandomForest	0.05	1.25	1.3	44.73
OneHotEnCoding	Stacking	0.59	1.1	1.18	38.19
OneHotEnCoding	Voting	0.84	1.12	1.2	37.74

1. After Applying TfidfVectorizer with top 2000 words, CV Log Loss for LogisticRegression with Class Balancing = 0.9
2. After Applying TfidfVectorizer with top 2000 words, CV Log Loss for LogisticRegression without Class Balancing = 0.99
3. After Applying TfidfVectorizer with top 2000 words, CV Log Loss for LinearSVM = 0.98